

Deployment Steps

HCI

NetApp
July 31, 2020

This PDF was generated from https://docs.netapp.com/us-en/hci-solutions/anthos_reference_workflow_summary.html on July 31, 2020. Always check docs.netapp.com for the latest.



Table of Contents

Deployment Steps	1
Workflow Summary: NetApp HCI with Anthos	1
1. Configure Management Switches: NetApp HCI with Anthos	1
2. Configure Data Switches: NetApp HCI with Anthos	5
3. Deploy NetApp HCI with the NetApp Deployment Engine: NetApp HCI with Anthos	11
4. Configure the vCenter Server: NetApp HCI with Anthos	21
5. Deploy and Configure the F5 Big-IP Virtual Edition Appliance: NetApp HCI with Anthos	25
6. Complete Anthos Prerequisites: NetApp HCI with Anthos	29
7. Deploy the Anthos Admin Workstation: NetApp HCI with Anthos	35
8. Deploy the Admin and the First User Cluster: NetApp HCI with Anthos	44
9. Deploy Additional User Clusters: NetApp HCI with Anthos	56
10. Enable Access to the Cluster with the GKE Console: NetApp HCI with Anthos	60
11. Install and Configure NetApp Trident Storage Provisioner: NetApp HCI with Anthos	63

Deployment Steps

Workflow Summary: NetApp HCI with Anthos

This section provides detailed protocols for implementing the NetApp HCI solution for Anthos.

This deployment is divided into the following high-level tasks:

1. Configure management switches
2. Configure data switches
3. Deploy NetApp HCI with the NetApp Deployment Engine
4. Configure the vCenter Server
5. Deploy and configure the F5 Big-IP Virtual Edition appliance
6. Complete Anthos prerequisites
7. Deploy the Anthos admin workstation
8. Deploy the admin and the first user cluster
9. Deploy additional user clusters
10. Enable access to the cluster with the GKE console
11. Install and configure NetApp Trident storage provisioner

1. Configure Management Switches: NetApp HCI with Anthos

Cisco Nexus 3048 switches are used in this deployment procedure to provide 1Gbps connectivity for in and out-of-band management of the compute and storage nodes. These steps begin after the switches have been racked, powered, and put through the initial setup process. To configure the switches to provide management connectivity to the infrastructure, complete the following steps:

Enable Advanced Features for Cisco Nexus

Run the following commands on each Cisco Nexus 3048 switch to configure advanced features:

1. Enter configuration mode.

```
Switch-01# configure terminal
```

2. Enable VLAN functionality.

```
Switch-01(config)# feature interface-vlan
```

3. Enable LACP.

```
Switch-01(config)# feature lacp
```

4. Enable virtual port channels (vPCs).

```
Switch-01(config)# feature vpc
```

5. Set the global port-channel load-balancing configuration.

```
Switch-01(config)# port-channel load-balance src-dst ip-l4port
```

6. Perform the global spanning-tree configuration.

```
Switch-01(config)# spanning-tree port type network default
Switch-01(config)# spanning-tree port type edge bpduguard default
```

Configure Ports on the Switch for In-Band Management

1. Run the following commands to create VLANs for management purposes.

```
Switch-01(config)# vlan 2
Switch-01(config-vlan)# Name Native_VLAN
Switch-01(config-vlan)# vlan 16
Switch-01(config-vlan)# Name OOB_Network
Switch-01(config-vlan)# vlan 3480
Switch-01(config-vlan)# Name MGMT_Network
Switch-01(config-vlan)# exit
```

2. Configure the ports ETH1/29-32 as VLAN trunk ports that connect to management interfaces on each HCI storage node.

```
Switch-01(config)# int eth 1/29
Switch-01(config-if)# description HCI-STG-01 PortA
Switch-01(config-if)# switchport mode trunk
Switch-01(config-if)# switchport trunk native vlan 2
Switch-01(config-if)# switchport trunk allowed vlan 3480
Switch-01(config-if)# spanning tree port type edge trunk
Switch-01(config-if)# int eth 1/30
Switch-01(config-if)# description HCI-STG-02 PortA
Switch-01(config-if)# switchport mode trunk
Switch-01(config-if)# switchport trunk native vlan 2
Switch-01(config-if)# switchport trunk allowed vlan 3480
Switch-01(config-if)# spanning tree port type edge trunk
Switch-01(config-if)# int eth 1/31
Switch-01(config-if)# description HCI-STG-03 PortA
Switch-01(config-if)# switchport mode trunk
Switch-01(config-if)# switchport trunk native vlan 2
Switch-01(config-if)# switchport trunk allowed vlan 3480
Switch-01(config-if)# spanning tree port type edge trunk
Switch-01(config-if)# int eth 1/32
Switch-01(config-if)# description HCI-STG-04 PortA
Switch-01(config-if)# switchport mode trunk
Switch-01(config-if)# switchport trunk native vlan 2
Switch-01(config-if)# switchport trunk allowed vlan 3480
Switch-01(config-if)# spanning tree port type edge trunk
Switch-01(config-if)# exit
```

Configure Ports on the Switch for Out-of-Band Management

1. Run the following commands to configure the ports for cabling the IPMI interfaces on each HCI node.

```
Switch-01(config)# int eth 1/13
Switch-01(config-if)# description HCI-CMP-01 IPMI
Switch-01(config-if)# switchport mode access
Switch-01(config-if)# switchport access vlan 16
Switch-01(config-if)# spanning-tree port type edge
Switch-01(config-if)# int eth 1/14
Switch-01(config-if)# description HCI-STG-01 IPMI
Switch-01(config-if)# switchport mode access
Switch-01(config-if)# switchport access vlan 16
Switch-01(config-if)# spanning-tree port type edge
Switch-01(config-if)# int eth 1/15
Switch-01(config-if)# description HCI-STG-03 IPMI
Switch-01(config-if)# switchport mode access
Switch-01(config-if)# switchport access vlan 16
Switch-01(config-if)# spanning-tree port type edge
Switch-01(config-if)# exit
```



In the validated configuration, we cabled odd-node IPMI interfaces to Switch-01, and even-node IPMI interfaces to Switch-02.

Create a vPC Domain to Ensure Fault Tolerance

1. Activate the ports used for the vPC peer-link between the two switches.

```
Switch-01(config)# int eth 1/1
Switch-01(config-if)# description vPC peer-link Switch-02 1/1
Switch-01(config-if)# int eth 1/2
Switch-01(config-if)# description vPC peer-link Switch-02 1/2
Switch-01(config-if)# exit
```

2. Perform the vPC global configuration.

```
Switch-01(config)# vpc domain 1
Switch-01(config-vpc-domain)# role priority 10
Switch-01(config-vpc-domain)# peer-keepalive destination <switch-02_mgmt_address>
source <switch-01_mgmt_address> vrf managment
Switch-01(config-vpc-domain)# peer-gateway
Switch-01(config-vpc-domain)# auto recovery
Switch-01(config-vpc-domain)# ip arp synchronize
Switch-01(config-vpc-domain)# int eth 1/1-2
Switch-01(config-vpc-domain)# channel-group 10 mode active
Switch-01(config-vpc-domain)# int Po10
Switch-01(config-if)# description vPC peer-link
Switch-01(config-if)# switchport mode trunk
Switch-01(config-if)# switchport trunk native vlan 2
Switch-01(config-if)# switchport trunk allowed vlan 16,3480
Switch-01(config-if)# spanning-tree port type network
Switch-01(config-if)# vpc peer-link
Switch-01(config-if)# exit
```

2. Configure Data Switches: NetApp HCI with Anthos

Mellanox SN2010 switches provide 25Gbps connectivity for the data plane of the compute and storage nodes.

To configure the switches to provide data connectivity to the infrastructure, complete the following steps:

Create MLAG Cluster to Ensure Fault Tolerance

1. Run the following commands on each Mellanox SN210 switch for general configuration:

1. Enter configuration mode.

```
Switch-01 enable
Switch-01 configure terminal
```

2. Enable the LACP required for the Inter-Peer Link (IPL).

```
Switch-01 (config) # lacp
```

3. Enable the Link Layer Discovery Protocol (LLDP).

```
Switch-01 (config) # lldp
```

4. Enable IP routing.

```
Switch-01 (config) # ip routing
```

5. Enable the MLAG protocol

```
Switch-01 (config) # protocol mlag
```

6. Enable global QoS.

```
Switch-01 (config) # dcb priority-flow-control enable force
```

2. For MLAG to function, the switches must be made peers to each other through an IPL. This should consist of two or more physical links for redundancy. The MTU for the IPL is set for jumbo frames (9216), and all VLANs are enabled by default. Run the following commands on each switch in the domain:

1. Create port channel 10 for the IPL.

```
Switch-01 (config) # interface port-channel 10
Switch-01 (config interface port-channel 10) # description IPL
Switch-01 (config interface port-channel 10) # exit
```

2. Add interfaces ETH 1/20 and 1/22 to the port channel.

```
Switch-01 (config) # interface ethernet 1/20 channel-group 10 mode active
Switch-01 (config) # interface ethernet 1/20 description ISL-SWB_01
Switch-01 (config) # interface ethernet 1/22 channel-group 10 mode active
Switch-01 (config) # interface ethernet 1/22 description ISL-SWB_02
```

3. Create a VLAN outside of the standard range dedicated to IPL traffic.

```
Switch-01 (config) # vlan 4000
Switch-01 (config vlan 4000) # name IPL VLAN
Switch-01 (config vlan 4000) # exit
```

4. Define the port channel as the IPL.

```
Switch-01 (config) # interface port-channel 10 ipl 1
Switch-01 (config) # interface port-channel 10 dcb priority-flow-control mode on
force
```

5. Set an IP for each IPL member (non-routable; it is not advertised outside of the switch).

```
Switch-01 (config) # interface vlan 4000
Switch-01 (config vlan 4000) # ip address 10.0.0.1 255.255.255.0
Switch-01 (config vlan 4000) # ipl 1 peer-address 10.0.0.2
Switch-01 (config vlan 4000) # exit
```

3. Create a unique MLAG domain name for the two switches and assign a MLAG virtual IP (VIP). This IP is used for keep-alive heartbeat messages between the two switches. Run these commands on each switch in the domain:

1. Create the MLAG domain and set the IP address and subnet.

```
Switch-01 (config) # mlag-vip MLAG-VIP-DOM ip a.b.c.d /24 force
```

2. Create a virtual MAC address for the system MLAG.

```
Switch-01 (config) # mlag system-mac AA:BB:CC:DD:EE:FF
```

3. Configure the MLAG domain so that it is active globally.

```
Switch-01 (config) # no mlag shutdown
```



The IP used for the MLAG VIP must be in the same subnet as the switch management network (mgmt0).



The MAC address used can be any unicast MAC address and must be set to the same value on both switches in the MLAG domain.

Configure Ports to Connect to Storage and Compute Hosts

1. Create each of the VLANs needed to support the services for NetApp HCI. Run these commands on each switch in the domain:

1. Create VLANs.

```
Switch-01 (config) # vlan 1172
Switch-01 (config vlan 1172) exit
Switch-01 (config) # vlan 3480-3482
Switch-01 (config vlan 3480-3482) exit
```

2. Create names for each VLAN for easier accounting.

```
Switch-01 (config) # vlan 1172 name VM_Network
Switch-01 (config) # vlan 3480 name MGMT_Network
Switch-01 (config) # vlan 3481 name Storage_Network
Switch-01 (config) # vlan 3482 name vMotion_Network
+
```

2. Create Hybrid VLAN ports on ports ETH1/9-10 so that you can tag the appropriate VLANs for the NetApp HCI compute nodes.

1. Select the ports you want to work with.

```
Switch-01 (config) # interface ethernet 1/9-1/10
```

2. Set the MTU for each port.

```
Switch-01 (config interface ethernet 1/9-1/10) # mtu 9216 force
```

3. Modify spanning-tree settings for each port.

```
Switch-01 (config interface ethernet 1/9-1/10) # spanning-tree bpdufilter enable
Switch-01 (config interface ethernet 1/9-1/10) # spanning-tree port type edge
Switch-01 (config interface ethernet 1/9-1/10) # spanning-tree bpduguard enable
```

4. Set the switchport mode to hybrid.

```
Switch-01 (config interface ethernet 1/9-1/10 ) # switchport mode hybrid
Switch-01 (config interface ethernet 1/9-1/10 ) # exit
```

5. Create descriptions for each port being modified.

```
Switch-01 (config) # interface ethernet 1/9 description HCI-CMP-01 PortD
Switch-01 (config) # interface ethernet 1/10 description HCI-CMP-02 PortD
```

6. Tag the appropriate VLANs for the NetApp HCI environment.

```
Switch-01 (config) # interface ethernet 1/9 switchport hybrid allowed-vlan add 1172
Switch-01 (config) # interface ethernet 1/9 switchport hybrid allowed-vlan add
3480-3482
Switch-01 (config) # interface ethernet 1/10 switchport hybrid allowed-vlan add
1172
Switch-01 (config) # interface ethernet 1/10 switchport hybrid allowed-vlan add
3480-3482
```

3. Create MLAG interfaces and hybrid VLAN ports on ports ETH1/5-8 so that you can distribute connectivity between the switches and tag the appropriate VLANs for the NetApp HCI storage nodes.

1. Select the ports that you want to work with.

```
Switch-01 (config) # interface ethernet 1/5-1/8
```

2. Set the MTU for each port.

```
Switch-01 (config interface ethernet 1/5-1/8) # mtu 9216 force
```

3. Modify spanning tree settings for each port.

```
Switch-01 (config interface ethernet 1/5-1/8) # spanning-tree bpdufilter enable
Switch-01 (config interface ethernet 1/5-1/8) # spanning-tree port type edge
Switch-01 (config interface ethernet 1/5-1/8) # spanning-tree bpduguard enable
```

4. Set the switchport mode to hybrid.

```
Switch-01 (config interface ethernet 1/5-1/8 ) # switchport mode hybrid
Switch-01 (config interface ethernet 1/5-1/8 ) # exit
```

5. Create descriptions for each port being modified.

```
Switch-01 (config) # interface ethernet 1/5 description HCI-STG-01 PortD
Switch-01 (config) # interface ethernet 1/6 description HCI-STG-02 PortD
Switch-01 (config) # interface ethernet 1/7 description HCI-STG-03 PortD
Switch-01 (config) # interface ethernet 1/8 description HCI-STG-04 PortD
```

6. Create and configure the MLAG port channels.

```
Switch-01 (config) # interface mlag-port-channel 115-118
Switch-01 (config interface mlag-port-channel 115-118) # exit
Switch-01 (config) # interface mlag-port-channel 115-118 no shutdown
Switch-01 (config) # interface mlag-port-channel 115-118 mtu 9216 force
Switch-01 (config) # interface mlag-port-channel 115-118 lacp-individual enable
force
Switch-01 (config) # interface ethernet 1/5-1/8 lacp port-priority 10
Switch-01 (config) # interface ethernet 1/5-1/8 lacp rate fast
Switch-01 (config) # interface ethernet 1/5 mlag-channel-group 115 mode active
Switch-01 (config) # interface ethernet 1/6 mlag-channel-group 116 mode active
Switch-01 (config) # interface ethernet 1/7 mlag-channel-group 117 mode active
Switch-01 (config) # interface ethernet 1/8 mlag-channel-group 118 mode active
```

7. Tag the appropriate VLANs for the storage environment.

```
Switch-01 (config) # interface mlag-port-channel 115-118 switchport mode hybrid
Switch-01 (config) # interface mlag-port-channel 115 switchport hybrid allowed-vlan
add 1172 Switch-01 (config) # interface mlag-port-channel 116 switchport hybrid
allowed-vlan add 1172
Switch-01 (config) # interface mlag-port-channel 117 switchport hybrid allowed-vlan
add 1172
Switch-01 (config) # interface mlag-port-channel 118 switchport hybrid allowed-vlan
add 1172
Switch-01 (config) # interface mlag-port-channel 115 switchport hybrid allowed-vlan
add 3481
Switch-01 (config) # interface mlag-port-channel 116 switchport hybrid allowed-vlan
add 3481
Switch-01 (config) # interface mlag-port-channel 117 switchport hybrid allowed-vlan
add 3481
Switch-01 (config) # interface mlag-port-channel 118 switchport hybrid allowed-vlan
add 3481
```



The configurations in this section must also be run on the second switch in the MLAG domain. NetApp recommends that the descriptions for each port are updated to reflect the device ports that are being cabled and configured on the other switch.

Create Uplink Ports for the Switches

1. Create an MLAG interface to provide uplinks to both Mellanox SN2010 switches from the core network.

```
Switch-01 (config) # interface mlag port-channel 101
Switch-01 (config interface mlag port-channel) # description Uplink CORE-SWITCH port
PORT
Switch-01 (config interface mlag port-channel) # exit
```

2. Configure the MLAG members.

```
Switch-01 (config) # interface ethernet 1/18 description Uplink to CORE-SWITCH port
PORT
Switch-01 (config) # interface ethernet 1/18 speed 10000 force
Switch-01 (config) # interface mlag-port-channel 101 mtu 9216 force
Switch-01 (config) # interface ethernet 1/18 mlag-channel-group 101 mode active
```

3. Set the switchport mode to hybrid and allow all VLANs from the core uplink switches.

```
Switch-01 (config) # interface mlag-port-channel switchport mode hybrid
Switch-01 (config) # interface mlag-port-channel switchport hybrid allowed-vlan all
```

4. Verify that the MLAG interface is up.

```
Switch-01 (config) # interface mlag-port-channel 101 no shutdown
Switch-01 (config) # exit
```

3. Deploy NetApp HCI with the NetApp Deployment Engine: NetApp HCI with Anthos

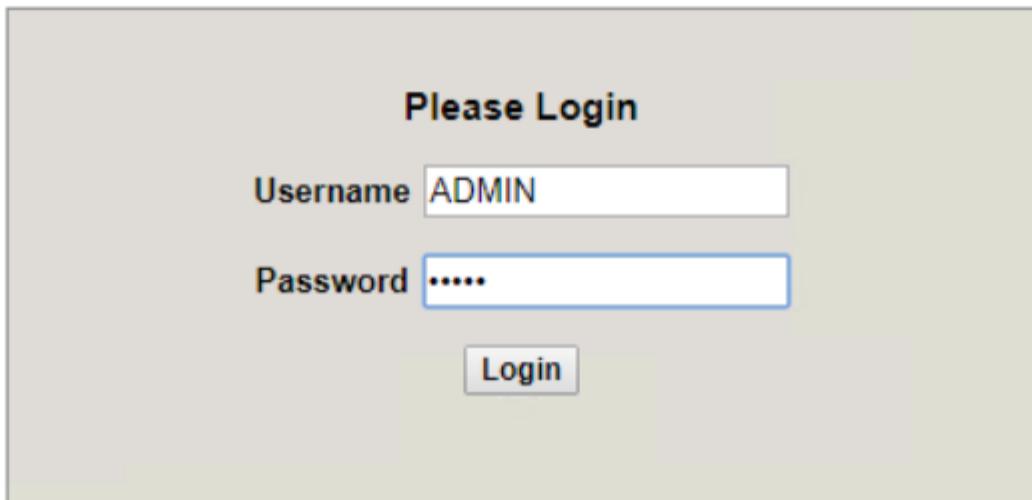
NDE delivers a simple and streamlined deployment experience for the NetApp HCI solution.

These steps begin after the nodes have been racked, and cabled, and the IPMI port has been configured on each node using the console.

A detailed guide to using NDE 1.6 to deploy your NetApp HCI system can be found [here](#).

To Deploy the NetApp HCI solution using NDE, complete the following steps:

1. Access the out-of-band management console for one of the storage nodes in the cluster and log in with the default credentials ADMIN/ADMIN.



The image shows a login interface for a NetApp system. At the top center, it says "Please Login". Below that, there is a "Username" field containing "ADMIN". Below the username field is a "Password" field containing "*****". At the bottom center is a "Login" button.

2. Click the Remote Console Preview image in the center of the screen to download a JNLP file launched by Java Web Start, which launches an interactive console to the system.
3. With the virtual console launched, a user can log in to the HCI Storage node, using the ADMIN/ADMIN username and password combination.
4. The Bond1G interface must have an IP, a netmask, and a gateway set statically; its VLAN set to 3480; and DNS servers defined for the environment.

```
Bond10G
  Method          : static
  Link Speed     : 50000
  IPv4 Address   :
  IPv4 Subnet Mask  :
    --->
  IPv4 Gateway Address :
    --->
  MTU            : 9000
    --->
  Bond Mode      : LACP  [ActivePassive, ALB, LACP]
    --->
  LACP Rate       : Fast   [Fast, Slow]
    --->
  Status          : UpAndRunning  [Down, Up, UpAndRunning]
    --->
  Virtual Network Tag :
    --->
  Routes          : Number of routes: 0.
    --->
```



Select an IP that is within the subnet you intend to use for in-band management, but not an IP you would like to use in production. NDE reconfigures the node with a production IP after initial access.



This task must only be performed on the first storage node. Afterward, the other nodes in the infrastructure are discovered by the Automatic Private IP Address (APIPA) addresses assigned to each storage interface when left unconfigured.

5. The Bond 10G interface must have its MTU setting changed to enable jumbo frames and its bond mode changed to LACP.

```

Bond10G
      Method          : static

      Link Speed     : 50000

      IPv4 Address    :

      IPv4 Subnet Mask   :
      --->

      IPv4 Gateway Address :
      --->

      MTU             : 9000
      --->

      Bond Mode       : LACP  [ActivePassive, ALB, LACP]
      --->

      LACP Rate        : Fast   [Fast, Slow]
      --->

      Status           : UpAndRunning  [Down, Up, UpAndRunning]
      --->

      Virtual Network Tag  :
      --->

      Routes          : Number of routes: 0.
      --->

```



Configure each of the four storage nodes in the NetApp HCI solution this way. The NDE process is then able to discover all the nodes in the solution and configure them. You do not need to modify the Bond10g interfaces on the two compute nodes.

6. After completion, open a web browser and visit the IP address you configured for the management port to start the NetApp HCI configuration with NDE.
7. On the Welcome to NetApp HCI page, click the Get Started button.
8. Check each associated box on the Prerequisites page and click Continue.
9. The next page presents End User Licenses for NetApp HCI and VMware vSphere. If you accept the terms, click I Accept at the end of each agreement and then click Continue.
10. Click Configure a New vSphere Deployment, select vSphere 6.5U2, and enter the Fully Qualified Domain Name (FQDN) of your vCenter Server. Then click Continue.

vSphere Configuration

You may elect to configure a new vSphere deployment or to join an existing vSphere deployment.

- Configure a new vSphere deployment
- Configure Using vSphere Version 6.7 Update 1
- Configure Using vSphere Version 6.5 Update 2
- Join and extend an existing vSphere deployment

If you have set up a DNS record for your new vCenter server, then configure your server using its fully qualified domain name and DNS server IP address:

- Configure Using a Fully Qualified Domain Name Best Practice!

vCenter Server Fully Qualified Domain Name

anthos-vc.cie.netapp.com



Note: The domain name must resolve to an unused IP address.

DNS Server IP Address

10.61.184.251



If you have not set up a DNS record for your new vCenter server, you may configure using an IP address that we define:

- Configure Using an IP Address ?

Note: Once defined, the IP address cannot be changed.

Back

Continue

11. NDE asks for the credentials to be used in the environment. This is used for VMware vSphere, the NetApp Element storage cluster, and the NetApp Mnode, which provides management functionality for the cluster. When you are finished, click Continue.

Credentials

Define the user name and password that will be used for the storage cluster, vCenter, and the management node.

User Name

admin



Password

.....



Password must contain:

- ✓ At least 8 characters
- ✓ No more than 20 characters
- ✓ 1 uppercase letter that is not the first character
- ✓ 1 lowercase letter
- ✓ 1 of the following special characters: !@#\$
- ✓ Allowed characters: A-Z a-z 0-9 !@#\$
- ✓ 1 number that is not the last character

Re-enter Password

.....



[Back](#)

[Continue](#)

12. NDE then prompts for the network topology used to cable the NetApp HCI environment. The validated solution in this document has been deployed using the 2 Cable Option for the compute nodes, and the 4 Cable Option for the storage nodes. Click Continue.

Network Topology

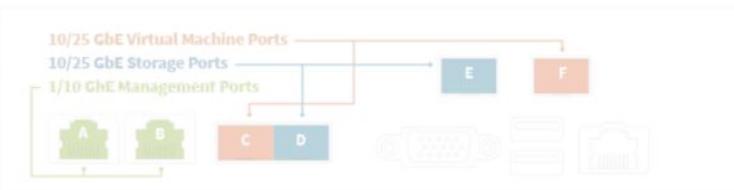
Select a compute node topology and a storage node topology appropriate for your hardware installation.

Compute Node Topology

6 Cable Option

The 6 cable option provides dedicated ports for management (2 x 1/10 GbE), virtual machines (2 x 10/25 GbE) and storage (2 x 10/25 GbE).

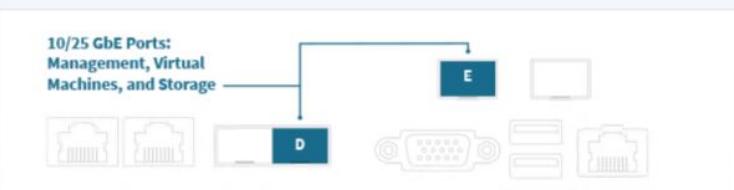
Use vSphere Distributed Switch? [?](#)



(H300E,H410C,H500E,H700E)

2 Cable Option

The 2 cable option provides shared management with ports for virtual machines and storage (2 x 10/25 GbE). The 2 cable option uses vSphere Distributed Switch. [?](#)

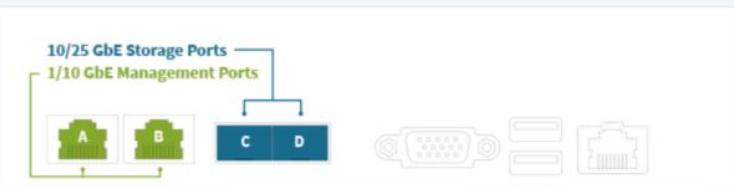


(H300E,H410C,H500E,H700E)

Storage Node Topology

4 Cable Option

The 4 cable option provides dedicated ports for management (2 x 1/10 GbE) and storage (2 x 10/25 GbE).



(H300S,H410S,H500S,H700S)

[Back](#)

[Continue](#)

13. The next page presented by NDE is the inventory of the environment as discovered by the APIPA addressed on the storage network. The storage node that is currently running NDE is already selected with a green check mark. Select the corresponding boxes to add additional nodes to the NetApp HCI environment. Click Continue.

Inventory

Verify the available nodes and select **at least 2 compute nodes and 4 storage nodes** to include in your installation.

[Refresh Inventory](#)

Compute Nodes

	Serial Number	Chassis Serial Number / Slot	Node Type	Software Version	Physical CPU Cores	Memory	1 GbE Ports	10 GbE Ports
<input checked="" type="checkbox"/>	HM17CS002729	002170990158 / B	H410C	1.6	8	384 GB	0 of 2 detected	2 of 4 detected
<input checked="" type="checkbox"/>	HM181S002024	002170990158 / A	H410C	1.6	8	384 GB	0 of 2 detected	2 of 4 detected

1 - 2 of 2 results

◀ ▶ 1 ▶ ▷

20 ▾

2 compute nodes selected

Storage Nodes

	Serial Number	Chassis Serial Number / Slot	Node Type	Raw Capacity	Element Version	Drive Count	1 GbE Ports	10 GbE Ports
<input checked="" type="checkbox"/>	221814003506	221814003436 / C	H500S	5.76 TB	11.3.1.5	6 of 6 detected	2 of 2 detected	2 of 2 detected
<input checked="" type="checkbox"/>	221818004613	221814003436 / D	H500S	5.76 TB	11.3.1.5	6 of 6 detected	2 of 2 detected	2 of 2 detected
<input checked="" type="checkbox"/> ?	221826005865	002170990158 / C	H500S	5.76 TB	11.3.1.5	6 of 6 detected	2 of 2 detected	2 of 2 detected
<input checked="" type="checkbox"/>	221826005866	002170990158 / D	H500S	5.76 TB	11.3.1.5	6 of 6 detected	2 of 2 detected	2 of 2 detected

1 - 4 of 4 results

◀ ▶ 1 ▶ ▷

20 ▾

4 storage nodes selected

[Back](#)

[Continue](#)



If there are any nodes missing from the inventory screen, wait a few minutes and click Refresh Inventory. If the node still fails to appear, additional investigation of environment networking may be required.

14. You must next configure the permanent network settings for the NetApp HCI deployment. The first page configures infrastructure services (DNS and NTP), vCenter networking, and Mnode networking.

Network Settings

Provide the network settings that will be used for your installation.

Live network validation is: **On**

Infrastructure Services

DNS Server IP Address 1	DNS Server IP Address 2 (Optional)
<input type="text" value="10.61.184.251"/>	<input type="text" value="10.61.184.252"/>
NTP Server Address 1	NTP Server Address 2 (Optional)
<input type="text" value="10.61.184.251"/>	<input type="text" value="10.61.184.252"/>

To save time, launch the easy form to enter fewer network settings.

vCenter Networking

VLAN ID	Subnet	Default Gateway	FQDN	IP Address
<input type="text" value="3480"/>	<input type="text" value="172.21.224.0/24"/>	<input type="text" value="172.21.224.1"/>	anthos-vc.cie.netapp.com	<input type="text" value="172.21.224.10"/>

Management Node Networking

Management Network		iSCSI Network
VLAN ID	Subnet	VLAN ID
<input type="text" value="3480"/>	<input type="text" value="172.21.224.0/24"/>	<input type="text" value="3481"/>
Subnet	Default Gateway	Subnet
<input type="text" value="172.21.224.0/24"/>	<input type="text" value="172.21.224.1"/>	<input type="text" value="172.21.225.0/24"/>
Default Gateway	Management IP Address	Storage (iSCSI) IP Address
<input type="text" value="anthos-mnode"/>	<input type="text" value="172.21.224.50"/>	<input type="text" value="172.21.225.50"/>

15. The next page allows you to configure each node in the environment. For the compute nodes, it allows you to configure the host name, management network, vMotion network, and storage network. For the storage nodes, name the storage cluster and configure the management and storage networks being used for each node. Click Continue.

Compute Node Networking

Management Network		vMotion Network		iSCSI A Network		iSCSI B Network	
VLAN ID	3480	VLAN ID	3482	VLAN ID	3481	VLAN ID	3481
Serial Number	Hostname	Management IP Address	vMotion IP Address	iSCSI A - IP Address	iSCSI B - IP Address		
HM17CS002729	Anthos-ESXi-01	172.21.224.11	172.21.226.11	172.21.225.11	172.21.225.111	172.21.225.111	172.21.225.111
HM181S002024	Anthos-ESXi-02	172.21.224.12	172.21.226.12	172.21.225.12	172.21.225.112	172.21.225.112	172.21.225.112

Storage Node Networking

Storage Cluster Name

Note: The storage cluster name cannot be changed after deployment.

Management Network		iSCSI Network	
VLAN ID	3480	VLAN ID	3481
Serial Number	Hostname	Management IP Address	Storage (iSCSI) IP Address
221814003506	Anthos-Store-01	172.21.224.21	172.21.225.21
221818004613	Anthos-Store-02	172.21.224.22	172.21.225.22
221826005865	Anthos-Store-03	172.21.224.23	172.21.225.23
221826005866	Anthos-Store-04	172.21.224.24	172.21.225.24

[Back](#)

Live network validation is: On [?](#)

[Continue](#)

- On the next page, review all the settings that have been defined for the environment by expanding each section, and, if necessary, click Edit to make corrections. There is also a check box on this page that enables or disables the Mnode from sending real-time health and diagnostics information to NetApp Active IQ®. If all the information is correct, click Start Deployment.



If you want to enable Active IQ, verify that your management network can reach the internet. If NDE is unable to reach Active IQ, the deployment can fail.

- A summary page appears along with a progress bar for each component of the NetApp HCI solution, as well as the overall solution. When complete, you are presented with an option to launch the vSphere client and begin working with your environment.

Your setup is complete.

[Launch vSphere Client](#)

Configure Network	Complete	✓
Set up NetApp Cluster	Complete	✓
Set up ESXi	Complete	✓
Set up vCenter	Complete	✓
Configure Management Node	Complete	✓
Finalize Configuration	Complete	✓

Overall Progress

100%

 [Export all setup information to CSV file](#)

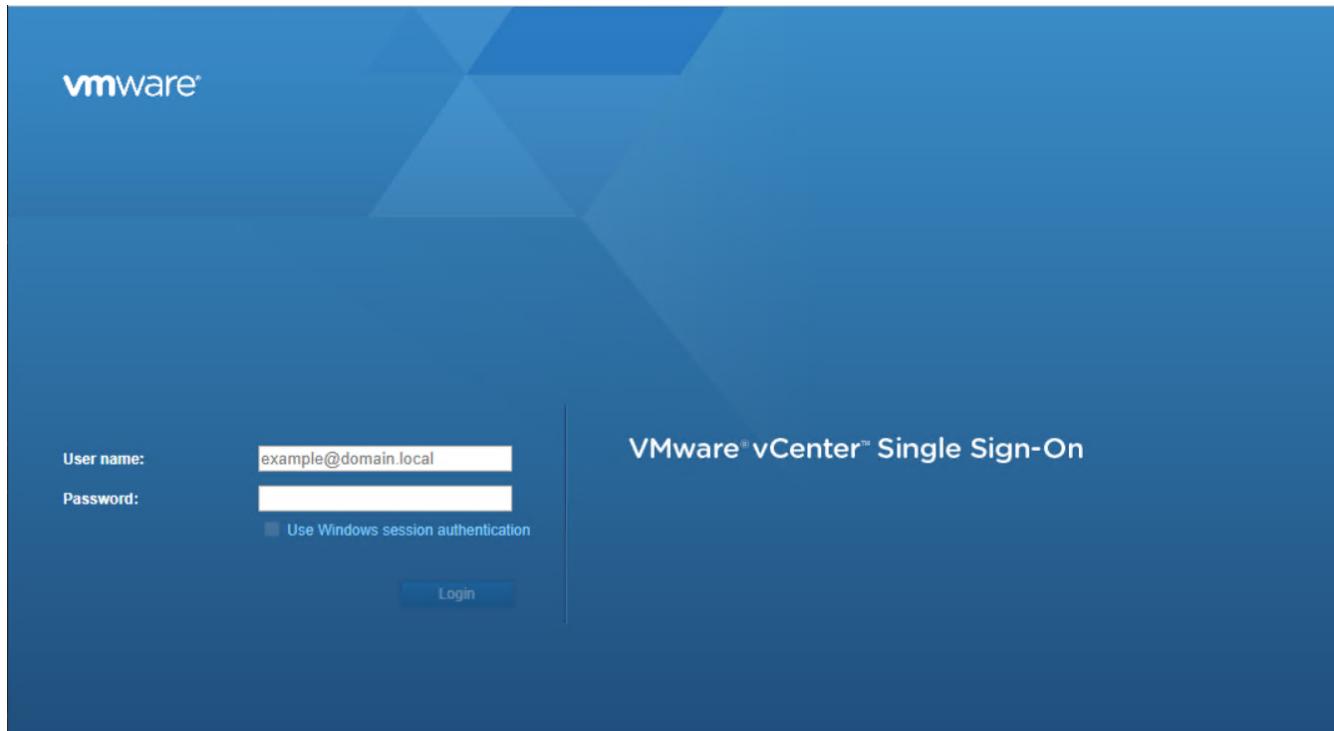
4. Configure the vCenter Server: NetApp HCI with Anthos

NDE deploys the solution with vCenter server and integrates the solution with the Element cluster by provisioning the Mnode VM and installing the NetApp Element Plug-in for vCenter.

After deployment, you must make a few modifications to the environment, including the creation of additional vDS portgroups, datastores, and resource groups for the deployment of the Anthos on VMware solution.

Complete the following steps to configure your vCenter Server for use:

1. Log into the VMware vCenter server using the [Administrator@vsphere.local](#) account and the password chosen for the admin user during NDE configuration.



2. Right-click **NetApp-HCI-Cluster-01** created by NDE and select the option to create a new resource pool. Name this pool **Infrastructure-Resource-Pool** and accept the defaults by clicking OK. This resource pool is used in a later configuration step.

New Resource Pool

NetApp-HCI-Cluster-01



Name	Infrastructure Resource		
CPU			
Shares	Normal	4000	
Reservation	0	MHz	▼
Max reservation: 54,128 MHz			
Reservation Type	<input checked="" type="checkbox"/> Expandable		
Limit	Unlimited	MHz	▼
Max limit: 58,128 MHz			
Memory			
Shares	Normal	163840	
Reservation	0	MB	▼
Max reservation: 751,064 MB			
Reservation Type	<input checked="" type="checkbox"/> Expandable		
Limit	Unlimited	MB	▼
Max limit: 756,820 MB			

CANCEL

OK



The reservations in this resource pool can be modified based on the resources available in the environment. NetApp HCI is deployed as an all-in-one solution. Therefore, NetApp recommends reserving the resources necessary to provide availability for the infrastructure services by placing them into this resource pool and adjusting the resources appropriately. Infrastructure services include vCenter Server, NetApp Mnode, and F5 Big-IP Load Balancer.

3. Repeat this step to create another resource pool for VMs deployed by Anthos. Name this pool Anthos-Resource-Pool and click the OK button to accept the default values. Adjust the resource availability based on the specific environment in which you are deploying the solution. This resource pool is used in a later deployment step.
4. To configure Element volumes to be used as vSphere datastores, click the dropdown menu and select NetApp Element Management from the list.
5. A Getting Started screen appears with details about your Element cluster.



6. Click Management, and the vSphere client presents a list of datastores. Click Create Datastore to create one datastore to host VMs and another to host ISOs for future guest installs.
7. Next click the Network menu item in the left panel. This displays a screen with information about the vDS deployed by NDE.
8. Several virtual port groups are defined by the initial configuration. NetApp recommends leaving these alone to support the infrastructure, and additional port groups should be created for user-deployed virtual guests. Right-click the NetApp HCI VDS 01 vDS in the left panel, and then select Distributed Port Group followed by the New Distributed Port Group option from the expanded menu.
9. Create a new distributed port group called **Management_Network**. Then click Next.
10. On the next screen, select the VLAN type as VLAN, and set the VLAN ID to 3480 for management purposes. Click Next, and, after reviewing the options on the summary page, click Next again to complete the creation of the distributed port group.
11. Repeat these steps to create distributed port groups for the **VM_Network** (VLAN 1172), as well as any other networks that might be used in the NetApp HCI environment.

 Additional networks can be defined to segment any additional deployed VMs. Examples of this use could be for a dedicated HA network for additional F5 Big-IP appliances if provisioned. Such configurations are in addition to the environment deployed in this validated solution and are considered out of scope for this NVA document.

5. Deploy and Configure the F5 Big-IP Virtual Edition Appliance: NetApp HCI with Anthos

Anthos enables native integration with F5 Big-IP load balancers to expose services from each pod to the world.

This solution makes use of the virtual appliance deployed in VMware vSphere as deployed by NDE. Networking for the F5 Big-IP virtual appliance can be configured in a two-armed or three-armed configuration based on your network environment. The deployment in this document is based on the two-armed configuration. Additional details on configuring the virtual appliance for use with Anthos can be found [here](#).

To deploy the F5 Big-IP Virtual Edition appliance, complete the following steps:

1. Download the virtual application Open Virtual Appliance (OVA) file from F5 [here](#).



To download the appliance, a user must register with F5. They provide a 30-day demo license for the Big-IP Virtual Edition Load Balancer. NetApp recommends a permanent 10Gbps license for the production deployment of an appliance.

2. Right-click the Infrastructure Resource Pool and select Deploy OVF Template. A wizard launches that allows you to select the OVA file that you just downloaded in Step 1. Click Next.

Deploy OVF Template

1 Select an OVF template

- 2 Select a name and folder
- 3 Select a compute resource
- 4 Review details
- 5 Select storage
- 6 Ready to complete

Select an OVF template

Select an OVF template from remote URL or local file system

Enter a URL to download and install the OVF package from the Internet, or browse to a location accessible from your computer, such as a local hard drive, a network share, or a CD/DVD drive.

URL

http | https://remoteserver-address/filetodeploy.ovf | .ova

Local file

BIGIP-15.0.1-0....ALL-vmware.ova

3. Click Next to continue through each step and accept the default values for each screen presented until you reach the storage selection screen. Select the VM_Datastore that was created earlier, and then click Next.
4. The next screen presented by the wizard allows you to customize the virtual networks for use in the environment. Select VM_Network for the External field and select Management_Network for the Management field. Internal and HA are used for advanced configurations for the F5 Big-IP appliance and are not configured. These parameters can be left alone, or they can be configured to connect to non-infrastructure, distributed port groups. Click Next.
5. Review the summary screen for the appliance, and, if all the information is correct, click Finish to start the deployment.
6. After the virtual appliance is deployed, right-click it and power it up. It should receive a DHCP address on the management network. The appliance is Linux-based, and it has VMware Tools deployed, so you can view the DHCP address it receives in the vSphere client.
7. Open a web browser and connect to the appliance at the IP address from the previous step. The

default login is admin/admin, and, after the first login, the appliance immediately prompts you to change the admin password. It then returns you to a screen where you must log in with the new credentials.

The screenshot shows the F5 BIG-IP Configuration Utility login interface. On the left, there are input fields for Hostname (bigip1), IP Address (172.21.224.20), Username (admin), and Password (*****). Below these fields is a 'Log in' button. On the right, there is a message area with the following text:
Welcome to the BIG-IP Configuration Utility.
Log in with your username and password using the fields on the left.

(c) Copyright 1996-2019, F5 Networks, Inc., Seattle, Washington. All rights reserved.
[F5 Networks, Inc. Legal Notices](#)

8. The first screen prompts the user to complete the Setup Utility. Begin the utility by clicking Next.
9. The next screen prompts for activation of the license for the appliance. Click Activate to begin. When prompted on the next page, paste either the 30-day evaluation license key you received when you registered for the download or the permanent license you acquired when you purchased the appliance. Click Next.
-  For the device to perform activation, the network defined on the management interface must be able to reach the internet.
10. On the next screen, the End User License Agreement (EULA) is presented. If the terms in the license are acceptable, click Accept.
11. The next screen counts the elapsed time as it verifies the configuration changes that have been made so far. Click Continue to resume with the initial configuration.
12. The Configuration Change window closes, and the Setup Utility displays the Resource Provisioning

menu. This window lists the features that are currently licensed and the current resource allocations for the virtual appliance and each running service.

13. Clicking the Platform menu option on the left enables additional modification of the platform. Modifications include setting the management IP address configured with DHCP, setting the host name and the time zone the appliance is installed in, and securing the appliance from SSH accessibility.
14. Next click the Network menu, which enables you to configure standard networking features. Click Next to begin the Standard Network Configuration wizard.
15. The first page of the wizard configures redundancy; leave the defaults and click Next. The next page enables you to configure an internal interface on the load balancer. Interface 1.1 maps to the vmnic labeled Internal in the OVF deployment wizard.

[Big-IP Configuration] | *big-IP_config_8.png*



The spaces in this page for Self IP Address, Netmask, and Floating IP address can be filled with a non-routable IP for use as a placeholder. They can also be filled with an internal network that has been configured as a distributed port group for virtual guests if you are deploying the three-armed configuration. They must be completed to continue with the wizard.

16. The next page enables you to configure an external network that is used to map services to the pods deployed in Kubernetes. Select a static IP from the VM_Network range, the appropriate subnet mask, and a floating IP from that same range. Interface 1.2 maps to the vmnic labeled External in the OVF deployment wizard.

[Big-IP Configuration] | *big-IP_config_9.png*

17. On the next page, you can configure an internal-HA network if you are deploying multiple virtual appliances in the environment. To proceed, you must fill the Self-IP Address and the Netmask fields, and you must select interface 1.3 as the VLAN Interface, which maps to the HA network defined by the OVF template wizard.
18. The next page enables you to configure the NTP servers. Then click Next to continue to the DNS setup. The DNS servers and domain search list should already be populated by the DHCP server. Click Next to accept the defaults and continue.
19. For the remainder of the wizard, click Next to continue through the advanced peering setup, the configuration of which is beyond the scope of this document. Then click Finish to exit the wizard.
20. Create individual partitions for the Anthos admin cluster and each user cluster deployed in the environment. Click System in the menu on the left, navigate to Users, and click Partition List.

21. The displayed screen only shows the current common partition. Click Create on the right to create the first additional partition, and name it **Anthos-Admin**. Then click Repeat, and name the partition **Anthos-Cluster1**, and click the Repeat button again to name the next partition **Anthos-Cluster2**. Finally click Finished to complete the wizard. The Partition list screen returns with all the partitions now listed.

6. Complete Anthos Prerequisites: NetApp HCI with Anthos

Now that the physical environment is set up, you can begin Anthos deployment. This starts with several prerequisites that you must meet to deploy the solution and access it afterward. Each of these steps are discussed in depth in the [Anthos GKE On-Prem Guide](#).

To prepare your environment for the deployment of Anthos on VMware, complete the following steps:

1. Create a Google Cloud project following the instructions available [here](#).



Your organization might already have a project in place intended for this purpose. Check with your cloud administration team to see if a project exists and is already configured for access to Anthos on VMware. All projects intended for use with Anthos must be whitelisted by Google. This includes the primary user account, additional team members, and the access service account created in a later step.

2. Create a deployment workstation from which to manage the installation of Anthos on VMware. The deployment workstation can be Linux, MacOS, or Windows. For the purposes of this validated deployment, Red Hat Enterprise Linux 7 was used.



This workstation can be hosted either internal or external to the NetApp HCI deployment. The only requirement is that it must be able to successfully communicate with the deployed VMware vCenter Server and the internet to function correctly.

3. Install [Google Cloud SDK](#) for interactions with Google Cloud. It can be downloaded as an archive of binaries for manual install or installed by either the apt-get (Ubuntu/Debian) or yum (RHEL) package managers.

```
[user@rhel7 ~]$ sudo yum install google-cloud-sdk
Failed to set locale, defaulting to C
Loaded plugins: langpacks, product-id, search-disabled-repos, subscription-manager
Resolving Dependencies
--> Running transaction check
--> Package google-cloud-sdk.noarch 0:270.0.0-1 will be installed
--> Finished Dependency Resolution
```

Dependencies Resolved

```
=====
=====
Package           Arch      Version       Repository
Size
=====
=====
Installing:
google-cloud-sdk      noarch    270.0.0-1   google-cloud-sdk
36 M
```

Transaction Summary

```
=====
=====
Install 1 Package
```

Total download size: 36 M

Installed size: 174 M

Is this ok [y/d/N]: y

Downloading packages:

```
6d81c821884ae40244c746f6044fc1bcd801143a0d9c8da06767036b8d090a24-google-cloud-sdk-
270.0.0-1.noar | 36 MB  00:00:00
```

Running transaction check

Running transaction test

Transaction test succeeded

Running transaction

```
  Installing : google-cloud-sdk-270.0.0-1.noarch
1/1
```

```
  Verifying  : google-cloud-sdk-270.0.0-1.noarch
1/1
```

Installed:

```
  google-cloud-sdk.noarch 0:270.0.0-1
```

Complete!



The gcloud binary must be at least version 265.0.0. You can update a manual install with a gcloud components update. However, if SDK was installed by a package manager, future updates must also be performed using that same package manager.

4. Install [govc](#), the CLI for VMware vSphere. Installing govc allows you to interact directly with the management of VMware vCenter. Govc is available as a pre-packaged binary available in a gzip format for download. For installation, a user must download the gzip archive, unzip, and copy the resulting binary to a local path directory such as /usr/local/bin.

```
[user@rhel7 ~]$ wget
https://github.com/vmware/govmomi/releases/download/v0.20.0/govc_linux_amd64.gz
--2019-11-06 09:06:10--
https://github.com/vmware/govmomi/releases/download/v0.20.0/govc_linux_amd64.gz
Resolving github.com (github.com)... 13.250.177.223
Connecting to github.com (github.com)|13.250.177.223|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 7650188 (7.3M) [application/octet-stream]
Saving to: 'govc_linux_amd64.gz.1'

100%[=====] 7,650,188  --.-K/s   in 0.1s

2019-11-06 09:06:12 (76.4 MB/s) - 'govc_linux_amd64.gz' saved [7650188/7650188]

[user@rhel7 ~]$ gunzip govc_linux_amd64.gz
[user@rhel7 ~]$ sudo cp govc_linux_amd64 /usr/local/bin/govc
[user@rhel7 ~]$ which govc
/usr/local/bin/govc
```

5. Install [Hashicorp Terraform](#). Terraform enables the automated deployment of VMs in a VMware vSphere environment and is used to deploy the Anthos admin workstation in a later step. Like the govc install, Terraform can be downloaded as a zip archive and unzipped, and the resulting binary can be copied to a directory in the local user's path.

```
[user@rhel7 ~]$ wget https://releases.hashicorp.com/terraform/0.12.13/terraform_0.12.13_linux_amd64.zip --2019-11-06 09:13:40-- https://releases.hashicorp.com/terraform/0.12.13/terraform_0.12.13_linux_amd64.zip Resolving releases.hashicorp.com (releases.hashicorp.com)... 151.101.193.183, 151.101.1.183, 151.101.65.183, ... Connecting to releases.hashicorp.com (releases.hashicorp.com)|151.101.193.183|:443... connected. HTTP request sent, awaiting response... 200 OK Length: 16341595 (16M) [application/zip] Saving to: 'terraform_0.12.13_linux_amd64.zip.1'

100%[=====>] 16,341,595  --.-K/s  in 0.1s

2019-11-06 09:13:40 (141 MB/s) - 'terraform_0.12.13_linux_amd64.zip.1' saved [16341595/16341595]

[user@rhel7 ~]$ unzip terraform_0.12.13_linux_amd64.zip
Archive:  terraform_0.12.13_linux_amd64.zip
  inflating: terraform
[user@rhel7 ~]$ cp terraform /usr/local/bin
[user@rhel7 ~]$ which terraform
/usr/local/bin/terraform
```

6. With the workstation configured, log in to Google Cloud with your credentials. To do so, enter the login command from the deployment workstation and retrieve a link that can be copied and pasted into a browser to allow interactive sign-in to Google services. After you have logged in, the web page presents a code that you can copy and paste back into the deployment workstation when prompted.

```
[user@rhel7 ~]$ gcloud auth login  
Go to the following link in your browser:
```

```
https://accounts.google.com/o/oauth2/auth?code_challenge=-  
7oPNSySHr_Sd2ZZ4K83koIeGTLVcdbjc8omr6zCbAI&prompt=select_account&code_challenge_method  
=S256&access_type=offline&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&response_ty  
pe=code&client_id=32655940559.apps.googleusercontent.com&scope=https%3A%3F%2Fwww.googl  
eapis.com%2Fauth%2Fuserinfo.email+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fcloud-  
platform+https%3A%6F%2Fwww.googleapis.com%2Fauth%2Fappengine.admin+https%3A%2F%2Fwww.g  
oogleapis.com%2Fauth%2Fcompute+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Faccounts.rau  
th
```

Enter verification code: 6/swGAh52VVgB-TRS5LVrSvP79ZdD1b9V60byUGqoY67a3zp9NPciIKsM

You are now logged in as [user@netapp.com].

Your current project is [anthos-dev]. You can change this setting by running:

```
$ gcloud config set project PROJECT_ID
```

7. Before you can install Anthos on VMware, you must create four service accounts, each with a specific purpose in interacting with Google Cloud. The following table lists the accounts and their purposes.

Table 1. Google Cloud Service Accounts

Account Name	Purpose
access-service-account	Used to download the Anthos binaries from Cloud Storage.
register-service-account	Used to register Anthos clusters to the Google Cloud console.
connect-service-account	Used to maintain the connection between user clusters and the Google Cloud.
stackdriver-service-account	Used to write logging and monitoring data to Stackdriver.



Each account is assigned an email address that references your approved Google Cloud project name. The examples below all list the project Anthos-Dev which was used during the NetApp validation. Make sure to substitute your appropriate project name in syntax examples where necessary.

```
[user@rhel7 ~]$ gcloud iam service-accounts create access-service-account
[user@rhel7 ~]$ gcloud iam service-accounts create register-service-account
[user@rhel7 ~]$ gcloud iam service-accounts create connect-service-account
[user@rhel7 ~]$ gcloud iam service-accounts create stackdriver-service-account
[user@rhel7 ~]$ gcloud iam service-accounts list
NAME          EMAIL
DISABLED
False          stackdriver-service-account@anthos-dev.iam.gserviceaccount.com
False          register-service-account@anthos-dev.iam.gserviceaccount.com
False          access-service-account@anthos-dev.iam.gserviceaccount.com
False          connect-service-account@anthos-dev.iam.gserviceaccount.com
False
```

8. Enable several APIs so that your environment can communicate with Google Cloud. The pods deployed in your clusters must be able to access <https://www.googleapis.com> and <https://gkeconnect.googleapis.com> to function as expected. Therefore, the VM_Network that the worker nodes are attached to must have internet access. To enable the necessary APIs, run the following command from the deployment workstation:

```
[user@rhel7 ~]$ gcloud services enable \
cloudresourcemanager.googleapis.com \
container.googleapis.com \
gkeconnect.googleapis.com \
gkehub.googleapis.com \
serviceusage.googleapis.com \
stackdriver.googleapis.com \
monitoring.googleapis.com \
logging.googleapis.com
```

9. The final step needed to prepare your environment to deploy Anthos is to limit certain privileges to your service accounts. You need the associated email address for each service account listed in Step 7.
 1. Using the register service account, assign the roles for `gkehub.admin` and `serviceusage.serviceUsageViewer`.

```
[user@rhel7 ~]$ gcloud projects add-iam-policy-binding anthos-dev \
--member "serviceAccount: register-service-account@anthos-
dev.iam.gserviceaccount.com" \
--role "roles/gkehub.admin"

[user@rhel7 ~]$ gcloud projects add-iam-policy-binding anthos-dev \
--member "serviceAccount: register-service-account@anthos-
dev.iam.gserviceaccount.com" \
--role "roles/serviceusage.serviceUsageViewer"
```

2. Using the connect service account, assign the roles for `gkehub.connect`.

```
[user@rhel7 ~]$ gcloud projects add-iam-policy-binding anthos-dev \
--member "serviceAccount: connect-service-account@anthos-
dev.iam.gserviceaccount.com" \
--role "roles/gkehub.connect"
```

3. With the stackdriver service account, assign the roles for `stackdriver.resourceMetadata.writer`, `logging.logWriter`, and `monitoring.metricWriter`.

```
[user@rhel7 ~]$ gcloud projects add-iam-policy-binding anthos-dev \
--member "serviceAccount: stackdriver-service-account@anthos-
dev.iam.gserviceaccount.com" \
--role "roles/stackdriver.resourceMetadata.writer"

[user@rhel7 ~]$ gcloud projects add-iam-policy-binding anthos-dev \
--member "serviceAccount: stackdriver-service-account@anthos-
dev.iam.gserviceaccount.com" \
--role "roles/logging.logWriter"

[user@rhel7 ~]$ gcloud projects add-iam-policy-binding anthos-dev \
--member "serviceAccount: stackdriver-service-account@anthos-
dev.iam.gserviceaccount.com" \
--role "roles/monitoring.metricWriter"
```

7. Deploy the Anthos Admin Workstation: NetApp HCI with Anthos

The admin workstation is a vSphere VM deployed within your NetApp HCI environment that is preinstalled with all the tools necessary to administer the Anthos on VMware solution. Follow the instructions in this section to download,

deploy, and configure the Anthos admin workstation.



The admin workstation image is packaged as an OVA file and is only available for download to those users with whitelisted access service accounts. If you are unable to download the OVA from cloud storage, contact your project administrator or open a support ticket with Google Cloud support.

To deploy the Anthos admin workstation, complete the following steps:

1. Download the latest version of the virtual appliance here. You must first set your account to the whitelisted access service account that has permission to download the OVA.

```
[user@rhel7 ~]$ gcloud config set account 'access-service-account@anthos-dev.iam.gserviceaccount.com'  
Updated property [core/account].  
[user@rhel7 ~]$ gsutil cp gs://gke-on-prem-release/admin-appliance/1.1.1-gke.2/gke-on-prem-admin-appliance-vsphere-1.1.1-gke.2.{ova,ova.sig} ~/  
Copying gs://gke-on-prem-release/admin-appliance/1.1.1-gke.2/gke-on-prem-admin-appliance-vsphere-1.1.1-gke.2.ova...  
==> NOTE: You are downloading one or more large file(s), which would run significantly faster if you enabled sliced object downloads. This feature is enabled by default but requires that compiled crcmod be installed (see "gsutil help crcmod").  
- [0 files][ 40.2 MiB/ 5.1 GiB]
```

2. When the OVA file is downloaded, use govc to load it into your vCenter. You need to set several environment variables for your current user that provide specific information about your VMware vSphere environment deployed on NetApp HCI.

```
[user@rhel7 ~]$ export GOVC_URL=https://anthos-vc.cie.netapp.com/sdk  
[user@rhel7 ~]$ export GOVC_USERNAME=administrator@vsphere.local  
[user@rhel7 ~]$ export GOVC_PASSWORD=vcpassword  
[user@rhel7 ~]$ export GOVC_DATASTORE=VM_Datastore  
[user@rhel7 ~]$ export GOVC_DATACENTER=NetApp-HCI-Datacenter-01  
[user@rhel7 ~]$ export GOVC_RESOURCE_POOL= Anthos-Resource-Pool  
[user@rhel7 ~]$ export GOVC_INSECURE=true
```

3. Because the NetApp HCI deployment uses a vDS, you need to create a custom options file that you pass as an additional argument to govc. Create a file called **options_specs**, and commit the following information to the file:

```

cat > options_spec <<EOF
{
  "DiskProvisioning": "thin",
  "MarkAsTemplate": true,
  "NetworkMapping": [
    {
      "Name": "VM Network",
      "Network": "VM_Network"
    }
  ]
}
EOF

```

- To import the OVA file as template in VMware vSphere, enter the following command referencing the downloaded virtual appliance file, as well as the options file that you just created. The console reports upload progress, and you can also verify the upload by browsing to the Recent Tasks window in VMware vSphere.

```

[user@rhel7 ~]$ govc import.ova -options options_spec .\gke-on-prem-admin-appliance-
vsphere-1.1.1-gke.2.ova
[06-11-19 10:56:39] Uploading gke-on-prem-admin-appliance-vsphere-1.1.1-gke.2.vmdk...
(23%, 26.8MiB/s)

```

Recent Tasks		Alarms	
Task Name	Target	Status	Initiator
Deploy OVF template	gke-on-prem-admin-a...	51%	VSPHERE.LOCAL\Adminis...
Check new notifications	anthos-vc.cle.netapp....	Completed	VMware vSphere Update ...
Refresh			
All			

- With the template uploaded, we can now use Terraform to deploy the Anthos admin workstation VM. Google provides a Terraform template and variable files for environments using DHCP as well those using static IP addresses. Create a directory for these Terraform files, and copy and paste the text for both the TF and TFPARS files available here into files created with a text editor.

```

[user@rhel7 ~]$ mkdir terraform
[user@rhel7 ~]$ cd terraform
[user@rhel7 ~]$ vi terraform.tfvars

```

- See the following [terraform.tfvars](#) file contents.

```
# vCenter Server username
vcenter_user = "administrator@vsphere.local"

# vCenter Server password
vcenter_password = "vcpassword"

# vCenter Server IP or hostname
vcenter_server = "https://anthos-vc.cie.netapp.com"

# Path in which the admin workstation's VM's public key should be saved
ssh_public_key_path = "~/.ssh/vsphere_workstation.pub"

# Hostname for the VM
vm_name = "admin-workstation"

# vSphere datastore to use for storage
datastore = "VM_Datastore"

# vSphere datacenter in which to create the VM
datacenter = "NetApp-HCI-Datacenter-01"

# vSphere cluster in which to create the VM
cluster = "NetApp-HCI-Cluster-01"

# vSphere resource pool in which to create VM, if you are using a non-default resource
pool
# If you are using the default resource pool, provide a value like "CLUSTER-
NAME/Resources"
resource_pool = "Anthos-Resource-Pool"

# vSphere network to use for the VM
network = "VM_Network"

# Number of CPUs for this VM. Recommended minimum 4.
num_cpus = 4

# Memory in MB for this VM. Recommended minimum 8192.
memory = 8192

# The VM template (OVA) to clone. Change the version if you imported a different
version of the OVA.
vm_template = "gke-on-prem-admin-appliance-vsphere-1.1.1-gke.2"
```

```
[user@rhel7 ~]$ vi terraform.tf
```

7. See the following `terraform.tf` file contents.

```
#####
##### VARIABLES #####
#####

# The following variables are declared in the accompanying TFVARS file

# vCenter Server username
variable "vcenter_user" { }

# vCenter Server password
variable "vcenter_password" { }

# vCenter Server address
variable "vcenter_server" { }

# Path in which the VM's public key should be saved
variable "ssh_public_key_path" { default = "~/.ssh/vsphere_workstation.pub" }

# vSphere network to use for the VM
variable "network" { default = "VM Network" }

# Hostname for the VM
variable "vm_name" { default = "vsphere-workstation" }

# vSphere datacenter in which to create the admin workstation VM
variable "datacenter" { }

# vSphere datastore to use for storage
variable "datastore" { }

# vSphere cluster in which to create the VM
variable "cluster" { }

# vSphere resource pool in which to create the VM
variable "resource_pool" { }

# Number of CPUs for this VM. Recommended minimum 4.
variable "num_cpus" { default = 4 }

# Memory in MB for this VM. Recommended minimum 8192.
variable "memory" { default = 8192 }

# The VM template (OVA) to clone
variable "vm_template" { }
```

```

#####
#####

provider "vsphere" {
  version      = "~> 1.5"
  user         = "${var.vcenter_user}"
  password     = "${var.vcenter_password}"
  vcenter_server = "${var.vcenter_server}"

  # if you have a self-signed cert
  allow_unverified_ssl = true
}

### vSphere Data ###

data "vsphere_datastore" "datastore" {
  name          = "${var.datastore}"
  datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_datacenter" "dc" {
  name = "${var.datacenter}"
}

data "vsphere_compute_cluster" "cluster" {
  name          = "${var.cluster}"
  datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_resource_pool" "pool" {
  name          = "${var.resource_pool}"
  datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_network" "network" {
  name          = "${var.network}"
  datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "vsphere_virtual_machine" "template_from_ovf" {
  name          = "${var.vm_template}"
  datacenter_id = "${data.vsphere_datacenter.dc.id}"
}

data "template_file" "dhcp_ip_config" {
  template = <<EOF
network:

```

```

version: 2
ethernets:
  ens192:
    dhcp4: true
EOF
}

data "template_file" "user_data" {
  template = <<EOF
#cloud-config
apt:
  primary:
    - arches: [default]
      uri: http://us-west1.gce.archive.ubuntu.com/ubuntu/
write_files:
  - path: /etc/netplan/99-dhcp.yaml
    permissions: '0644'
    encoding: base64
    content: |
      $$\{dhcp_ip_config\}
runcmd:
  - netplan apply
  - /var/lib/gke/guest-startup.sh
EOF
vars = {
  dhcp_ip_config = "${base64encode(data.template_file.dhcp_ip_config.rendered)}"
}
}

### vSphere Resources ###

resource "vsphere_virtual_machine" "vm" {
  name          = "${var.vm_name}"
  resource_pool_id = "${data.vsphere_resource_pool.pool.id}"
  datastore_id    = "${data.vsphere_datastore.datastore.id}"
  num_cpus        = "${var.num_cpus}"
  memory          = "${var.memory}"
  guest_id         = "${data.vsphere_virtual_machine.template_from_ovf.guest_id}"
  enable_disk_uuid = "true"
  scsi_type       = "${data.vsphere_virtual_machine.template_from_ovf.scsi_type}"
  network_interface {
    network_id   = "${data.vsphere_network.network.id}"
    adapter_type =
      "${data.vsphere_virtual_machine.template_from_ovf.network_interface_types[0]}"
  }
  wait_for_guest_net_timeout = 15
}

```

```

nested_hv_enabled = false
cpu_performance_counters_enabled = false

disk {
    label          = "disk0"
    size           = "${max(50,
data.vsphere_virtual_machine.template_from_ovf.disks.0.size)}"
    eagerly_scrub   =
"${data.vsphere_virtual_machine.template_from_ovf.disks.0.eagerly_scrub}"
    thin_provisioned =
"${data.vsphere_virtual_machine.template_from_ovf.disks.0.thin_provisioned}"
}

cdrom {
    client_device = true
}

vapp {
    properties = {
        hostname      = "${var.vm_name}"
        public-keys   = "${file(var.ssh_public_key_path)}"
        user-data     = "${base64encode(data.template_file.user_data.rendered)}"
    }
}

clone {
    template_uuid = "${data.vsphere_virtual_machine.template_from_ovf.id}"
}
}

output "ip_address" {
    value = "${vsphere_virtual_machine.vm.default_ip_address}"
}

```



Values specific to the deployed environment have been added to the `terraform.tfvars` file. However, you should not modify the `terraform.tf` file in any manner.

8. Create an SSH public/private keypair used to log in to the admin workstation after it is deployed. Name the public key so that it matches the variable that was assigned in the `terraform.tfvars` file.

```
[user@rhel7 ~]$ ssh-keygen -t rsa -f ~/.ssh/vsphere_workstation -N ""
Generating public/private rsa key pair.
Your identification has been saved in /home/user/.ssh/vsphere_workstation2.
Your public key has been saved in /home/user/.ssh/vsphere_workstation2.pub.
The key fingerprint is:
SHA256:qEk8G13LhwIKqf85ekHHZkIZduX2MkZUxGNEHvFT2vw user@rhel7
The key's randomart image is:
+---[RSA 2048]---+
|   oo.0+*B. . |
|   .o. o .+o = |
|   ... B.o.+ o |
| o ooo=B *   . |
|o ..==+ S .     E|
|. ...*       . |
|.   +.       . |
|.   o.       . |
|   o+o.       . |
+---[SHA256]---+
```

9. Navigate to the directory created to host the TF and TFPARS files. Within this directory, initialize Terraform and use it to launch the deployment of the admin workstation VM.

```

[user@rhel7 ~]$ cd terraform
[user@rhel7 terraform]$ ls
terraform.tf  terraform.tfvars
[user@rhel7 terraform]$ terraform init && terraform apply -auto-approve -input=false
Initializing the backend...
Initializing provider plugins...
The following providers do not have any version constraints in configuration,
so the latest version was installed.
To prevent automatic upgrades to new major versions that may contain breaking
changes, it is recommended to add version = "..." constraints to the
corresponding provider blocks in configuration, with the constraint strings
suggested below.
* provider.template: version = "~> 2.1"
Terraform has been successfully initialized!

data.template_file.dhcp_ip_config: Refreshing state...
data.template_file.user_data: Refreshing state...
data.vsphere_datacenter.dc: Refreshing state...
data.vsphere_resource_pool.pool: Refreshing state...
data.vsphere_datastore.datastore: Refreshing state...
data.vsphere_virtual_machine.template_from_ovf: Refreshing state...
data.vsphere_network.network: Refreshing state...
data.vsphere_compute_cluster.cluster: Refreshing state...
vsphere_virtual_machine.vm: Creating...
vsphere_virtual_machine.vm: Still creating... [10s elapsed]
vsphere_virtual_machine.vm: Still creating... [20s elapsed]
vsphere_virtual_machine.vm: Still creating... [30s elapsed]
vsphere_virtual_machine.vm: Still creating... [40s elapsed]
vsphere_virtual_machine.vm: Creation complete after 49s [id=42118cfa-d464-b815-f7b1-37cd85b2943a]

Warning: "vcenter_server": [DEPRECATED] This field has been renamed to vsphere_server.

Apply complete! Resources: 1 added, 0 changed, 0 destroyed.

Outputs:

ip_address = 10.63.172.21

```

8. Deploy the Admin and the First User Cluster: NetApp HCI with Anthos

All Kubernetes clusters deployed as a part of the Anthos solution are deployed from the Anthos admin workstation that you just created. A user logs into the

admin workstation using SSH, the public key created in a previous step, and the IP address provided at the end of the VM deployment. They can then begin creating their first clusters.



There are specific procedures for deploying clusters that use static IP addresses [here](#), and procedures for environments with DHCP can be found [here](#). In this guide, we use the second set of instructions.

To deploy the initial admin and first user cluster, complete the following steps:

1. Login to the admin workstation using the SSH public key and the IP address output at the end of Terraform deployment.

```
[user@rhel7 ~]$ ssh -i ~/.ssh/vsphere_workstation ubuntu@10.63.172.21
Welcome to Ubuntu 18.04.3 LTS (GNU/Linux 4.15.0-62-generic x86_64)
```

```
* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

This system has been minimized by removing packages and content that are not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.

```
36 packages can be updated.
23 updates are security updates.
```

```
To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.
```

```
ubuntu@admin-workstation:~$
```

2. Log in with the `gcloud auth` command as you did before from your deployment workstation by copying the URL into a web browser, signing into your Google account, and pasting the verification code back into the workstation.

```
ubuntu@admin-workstation:~$ gcloud auth login  
Go to the following link in your browser:
```

```
https://accounts.google.com/o/oauth2/auth?code_challenge=Q1F7H-  
CMUMuArasQD6AzHA0avKoLGfDqUyUgjFxf9ZI&prompt=select_account&code_challenge_method=S256  
&access_type=offline&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.0%3Aoob&response_type=co  
de&client_id=32555940559.apps.googleusercontent.com&scope=https%3A%2F%2Fwww.googleapis  
.com%2Fauth%2Fuserinfo.email+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fcloud-  
platform+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fappengine.admin+https%3A%2F%2Fwww.g  
oogleapis.com%2Fauth%2Fcompute+https%3A%2F%2Fwww.googleapis.com%2Fauth%2Faccounts.rau  
th
```

```
Enter verification code: 6/swFG_ZKZyd0eb1fuXsgtnBgYrIt1XwFQeB5ASrEqaaDfa09aio0bnNg
```

```
WARNING: `gcloud auth login` no longer writes application default credentials.
```

```
If you need to use ADC, see:
```

```
gcloud auth application-default --help
```

```
You are now logged in as [user@netapp.com].
```

```
Your current project is [None]. You can change this setting by running:
```

```
$ gcloud config set project PROJECT_ID
```

3. Set the project that you intend to deploy your clusters in for Anthos on VMWare, because the project is not set by default when you log in.

```
ubuntu@admin-workstation:~$ gcloud config set project anthos-dev  
Updated property [core/project].
```

4. Register gcloud as a Docker credential helper, which enables it to manage the credentials for Docker registries used for deployment. This way, the default credential store is not used for operations involving the credentials of the specified registries.

```
ubuntu@admin-workstation:~$ gcloud auth configure-docker
The following settings will be added to your Docker config file
located at [/home/ubuntu/.docker/config.json]:
{
  "credHelpers": {
    "gcr.io": "gcloud",
    "us.gcr.io": "gcloud",
    "eu.gcr.io": "gcloud",
    "asia.gcr.io": "gcloud",
    "staging-k8s.gcr.io": "gcloud",
    "marketplace.gcr.io": "gcloud"
  }
}
```

Do you want to continue (Y/n)? y

Docker configuration file updated.



By default, Anthos on VMware uses a pre-existing, Google-owned container image registry that requires no additional setup. If you choose to use a private Docker registry for deployment, then you must configure that registry separately based on instructions found [here](#). This step is beyond the scope of this deployment guide.

5. In the next step to deploy an admin cluster, create a private key file in the JSON format for each of the service accounts created in the prerequisites section.

```

ubuntu@admin-workstation:~$ gcloud iam service-accounts list
NAME          EMAIL
DISABLED
stackdriver-service-account@anthos-dev.iam.gserviceaccount.com
False
register-service-account@anthos-dev.iam.gserviceaccount.com
False
access-service-account@anthos-dev.iam.gserviceaccount.com
False
connect-service-account@anthos-dev.iam.gserviceaccount.com
False

ubuntu@admin-workstation:~$ gcloud iam service-accounts keys create access-key.json
--iam-account access-service-account@anthos-dev.iam.gserviceaccount.com
created key [8d5f8ce039dd98766e18a3c5ee6794912fb8d095] of type [json] as [access-
key.json] for [access-service-account@anthos-dev.iam.gserviceaccount.com]
ubuntu@admin-workstation:~$ gcloud iam service-accounts keys create register-key.json
--iam-account register-service-account@anthos-dev.iam.gserviceaccount.com
created key [f08b494c665321f83bcb8c8526ba21185b456a11] of type [json] as [register-
key.json] for [register-service-account@anthos-dev.iam.gserviceaccount.com]
ubuntu@admin-workstation:~$ gcloud iam service-accounts keys create connect-key.json
--iam-account connect-service-account@anthos-dev.iam.gserviceaccount.com
created key [c9640021ff6157d3df2a15db49f5c85b7b1495c2] of type [json] as [connect-
key.json] for [connect-service-account@anthos-dev.iam.gserviceaccount.com]
ubuntu@admin-workstation:~$ gcloud iam service-accounts keys create stackdriver-
key.json --iam-account stackdriver-service-account@anthos-dev.iam.gserviceaccount.com
created key [3c9427dfdef161d139ff998be896565c1df0b122] of type [json] as [stackdriver-
key.json] for [stackdriver-service-account@anthos-dev.iam.gserviceaccount.com]

```

- For the next step, you must use the access service key created in the previous step to activate the associated service account.

```

ubuntu@admin-workstation:~$ gcloud auth activate-service-account --key-file=access
-key.json
Activated service account credentials for: [access-service-account@anthos-
dev.iam.gserviceaccount.com]

```

- The deployment of the first clusters is performed using inputs from a config file generated by GKE. A generic config file can be created with no additional input, or an existing config file can be referenced to create additional clusters.

```
ubuntu@admin-workstation:~$ gkectl create-config  
ubuntu@admin-workstation:~$ ls  
access-key.json config.yaml connect-key.json register-key.json stackdriver-  
key.json
```

8. The **config.yaml** file created by running the previous command has several variables that must be customized for the current environment.

1. First, you must determine the full path location and the name of the current GKE bundle that is deployed into the environment. The file exists in the **/var/lib/gke/bundles** directory on the admin workstation.

```
ubuntu@admin-workstation:~$ ls /var/lib/gke/bundles  
gke-onprem-vsphere-1.1.1-gke.2-full.tgz gke-onprem-vsphere-1.1.1-gke.2.tgz
```

2. Next, you must get the fully recognized host name or IP address of our vCenter Server as displayed in its default SSL certificate. Connect to vSphere and dump the certificate contents into a file called **vcenter.pem**. Examining this file gives you the information that you need for the value of **Subject: CN** (common name).

```
ubuntu@admin-workstation:~$ true | openssl s_client -connect anthos-  
vc.cie.netapp.com:443 -showcerts 2>/dev/null | sed -ne '/-BEGIN/,/-END/p' >  
vcenter.pem  
ubuntu@admin-workstation:~$ openssl x509 -in vcenter.pem -text -noout | grep  
Subject:\ CN  
Subject: CN = anthos-vc.cie.netapp.com, C = US
```



If the value added to the **config.yaml** file does not match that of the CN found in the certificate, communication with the vCenter server fails.

9. With the information from those two commands and the generated **vcenter.pem** file, we can now edit the **config.yaml** file to prepare for deployment. Editing this file is very similar to the edits that you performed to the **terraform.tfvars** file to provide specifics regarding the VMware vCenter instance deployed in NetApp HCI.



When deploying the cluster, determine which IP addresses to use for the control plane and ingress VIPs for both the admin and user cluster. Also determine the compute and memory resources that must be reserved for each node deployed, because it is not possible to edit a cluster after it has been deployed.

```
# Absolute path to a GKE bundle on disk  
bundlepath: "/var/lib/gke/bundles/gke-onprem-vsphere-1.1.1-gke.2-full.tgz"
```

```

# Specify which vCenter resources to use for deployment
vcenter:
  # The credentials and address GKE should use to connect to vCenter
  credentials:
    address: "anthos-vc.cie.netapp.com"
    username: "administrator@vsphere.local"
    password: "vcpass"
  datacenter: "NetApp-HCI-Datacenter-01"
  datastore: "VM_Datastore"
  cluster: "NetApp-HCI-Cluster-01"
  network: "VM_Network"
  resourcepool: "Anthos Resource Pool"
  # Provide the name for the persistent disk to be used by the deployment (ending
  # in .vmdk). Any directory in the supplied path must be created before deployment.
  # Not required when adding additional user clusters
  datadisk: "anthos-admin-data-disk.vmdk"
  # Provide the path to vCenter CA certificate pub key for SSL verification
  cacertpath: "/home/ubuntu/vcenter.pem"

# Specify the proxy configuration.
proxy:
  # The URL of the proxy
  url: ""
  # The domains and IP addresses excluded from proxying
  noproxy: ""

# Specify admin cluster settings for a fresh GKE On-Prem deployment. Omit this section
# and use the --adminconfig flag when adding a new user cluster to an existing
# deployment
admincluster:
  # In-Cluster vCenter configuration
  vcenter:
    # If specified it overwrites the network field in global vcenter configuration
    network: ""
    # # The absolute or relative path to the yaml file to use for static IP allocation.
    # # Do not include if using DHCP
    # ipblockfilepath: ""
    # # Specify pre-defined nodeports if using "manual" load balancer mode
    # manuallbspec:
    #   ingresshttpnodeport: 32527
    #   ingresshttpsnodeport: 30139
    #   controlplanenodeport: 30968
    #   addonsnodeport: 31405
    # Specify the already-existing partition and credentials to use with F5
    bigip:
      # To re-use credentials across clusters we recommend using YAML node anchors.
      # See https://yaml.org/spec/1.2/spec.html#id2785586
      credentials:
        address: "172.21.224.22"
        username: "admin"

```

```
    password: "lbpass"
    partition: "Anthos-Admin-Part"
    # # Optionally specify a pool name if using SNAT
    # snatpoolname: ""
# The VIPs to use for load balancing
vips:
    # Used to connect to the Kubernetes API
    controlplanevip: "10.63.172.98"
    # Shared by all services for ingress traffic
    ingressvip: "10.63.172.99"
    # # Used for admin cluster addons (needed for multi cluster features). Must be the
same
    # # across clusters
    # addonsvip: ""
# The Kubernetes service CIDR range for the cluster. Must not overlap with the pod
# CIDR range
serviceiprange: 10.96.232.0/24
# The Kubernetes pod CIDR range for the cluster. Must not overlap with the service
# CIDR range
podiprange: 192.168.0.0/16
# Specify settings when deploying a new user cluster. Used both with a fresh
deployment
# or when adding a new cluster to an existing deployment.
usercluster:
    antiaffinitygroups:
        enabled: false
    # In-Cluster vCenter configuration
    vcenter:
        # If specified it overwrites the network field in global vcenter configuration
        network: ""
        # # The absolute or relative path to the yaml file to use for static IP allocation.
        # # Do not include if using DHCP
        # ipblockfilepath: ""
        # # Specify pre-defined nodeports if using "manual" load balancer mode
        # manuallbspec:
        #     ingresshttpnodeport: 30243
        #     ingresshttpsnodeport: 30879
        #     controlplanenodeport: 30562
        #     addonsnodeport: 0
        # Specify the already-existing partition and credentials to use with F5
        bigip:
            # To re-use credentials across clusters we recommend using YAML node anchors.
            # See https://yaml.org/spec/1.2/spec.html#id2785586
            credentials:
                address: "172.21.224.22"
                username: "admin"
                password: "lbpass"
            partition: "Anthos-Cluster01-Part"
```

```
# # Optionally specify a pool name if using SNAT
# snatpoolname: ""

# The VIPs to use for load balancing
vips:
    # Used to connect to the Kubernetes API
    controlplanevip: "10.63.172.105"
    # Shared by all services for ingress traffic
    ingressvip: "10.63.172.106"
    # # Used for admin cluster addons (needed for multi cluster features). Must be the
    same
        # # across clusters
        # addonsvip: ""

    # A unique name for this cluster
    clustername: "anthos-cluster01"

    # User cluster master nodes must have either 1 or 3 replicas
    masternode:
        cpus: 4
        memorymb: 8192
        # How many machines of this type to deploy
        replicas: 1

    # The number of worker nodes to deploy and their size. Min. 2 replicas
    workernode:
        cpus: 4
        memorymb: 8192
        # How many machines of this type to deploy
        replicas: 3

    # The Kubernetes service CIDR range for the cluster
    serviceiprange: 10.96.0.0/12
    # The Kubernetes pod CIDR range for the cluster
    podiprange: 192.168.0.0/16

    # # Uncomment this section to use OIDC authentication
    # oidc:
    #     issuerurl: ""
    #     kubectlredirecturl: ""
    #     clientid: ""
    #     clientsecret: ""
    #     username: ""
    #     usernameprefix: ""
    #     group: ""
    #     groupprefix: ""
    #     scopes: ""
    #     extraparams: ""
    #     # Set value to string "true" or "false"
    #     usehttpproxy: ""
    #     # # The absolute or relative path to the CA file (optional)
    #     # capath: ""
    #     # # Optionally provide an additional serving certificate for the API server
    #     sni:
```

```

#   certpath: ""
#   keypath: ""
# Which load balancer mode to use "Manual" or "Integrated"
lbmode: Integrated
# Specify which GCP project to connect your GKE clusters to
gkeconnect:
  projectid: "anthos-dev"
  # The absolute or relative path to the key file for a GCP service account used to
  # register the cluster
  registerserviceaccountkeypath: "/home/ubuntu/register-key.json"
  # The absolute or relative path to the key file for a GCP service account used by
  # the GKE connect agent
  agentserviceaccountkeypath: "/home/ubuntu/connect-key.json"
# Specify which GCP project to connect your logs and metrics to
stackdriver:
  projectid: "anthos-dev"
  # A GCP region where you would like to store logs and metrics for this cluster.
  clusterlocation: "us-east1"
  enablevpc: false
  # The absolute or relative path to the key file for a GCP service account used to
  # send logs and metrics from the cluster
  serviceaccountkeypath: "/home/ubuntu/stackdriver-key.json"
# # Optionally use a private Docker registry to host GKE images
# privaterRegistryconfig:
#   # Do not include the scheme with your registry address
#   credentials:
#     address: ""
#     username: ""
#     password: ""
#   # The absolute or relative path to the CA certificate for this registry
#   cacertpath: ""
# The absolute or relative path to the GCP service account key that will be used to
# pull GKE images
gcrkeypath: "/home/ubuntu/access-key.json"
# Configure kubernetes apiserver audit logging
cloudauditlogging:
  projectid: ""
  # A GCP region where you would like to store audit logs for this cluster.
  clusterlocation: ""
  # The absolute or relative path to the key file for a GCP service account used to
  # send audit logs from the cluster
  serviceaccountkeypath: ""

```

- Because spacing in YAML files can be very important, you can check the syntax of the config file by running the following command. If the command outputs any failures, be sure to examine the file and make any needed corrections.

```
ubuntu@admin-workstation:~$ gkectl check-config --config config.yaml
```

- Validation Category: Config Check
 - [SUCCESS] Config
- Validation Category: Docker Registry
 - [SUCCESS] gcr.io/gke-on-prem-release access
- Validation Category: vCenter
 - [SUCCESS] Credentials
 - [SUCCESS] Datacenter
 - [SUCCESS] Datastore
 - [SUCCESS] Data Disk
 - [SUCCESS] Resource Pool
 - [SUCCESS] Network
- Validation Category: F5 BIG-IP
 - [SUCCESS] Credentials
 - [SUCCESS] Partition
- Validation Category: Network Configuration
 - [SUCCESS] CIDR, VIP and static IP (availability and overlapping)
- Validation Category: VIPs
 - [SUCCESS] ping (availability)
- Validation Category: Node IPs
 - [SKIPPED] ping (availability): All specified clusters use DHCP.

Some validations FAILED or SKIPPED. Check report above.



Using DHCP skips the step to validate node IP availability. This is an expected behavior and deployment can continue.

11. Preparing the cluster for deployment and deploying the cluster are performed with two commands:
 1. The `gkectl prepare` command initializes the vSphere environment by uploading the node OS image, marking it as a template, and validating the build attestations for all container images.

```
ubuntu@admin-workstation:~$ gkectl prepare --config config.yaml
- Validation Category: Config Check
  - [SUCCESS] Config

- Validation Category: Docker Registry
  - [SUCCESS] gcr.io/gke-on-prem-release access

- Validation Category: vCenter
  - [SUCCESS] Credentials
  - [SUCCESS] Datacenter
  - [SUCCESS] Datastore
  - [SUCCESS] Data Disk
  - [SUCCESS] Resource Pool
  - [SUCCESS] Network
```

All validations SUCCEEDED.

```
Downloading OS image gke-on-prem-osimage-1.13.7-gke.20-20190816-8138298d96.ova...
DONE
```

Setting up OS image as a VM template in vSphere... DONE

2. The `gkectl create cluster` command deploys the cluster as depicted in the `config.yaml` file.

```
ubuntu@admin-workstation:~$ gkectl create cluster --config config.yaml
```

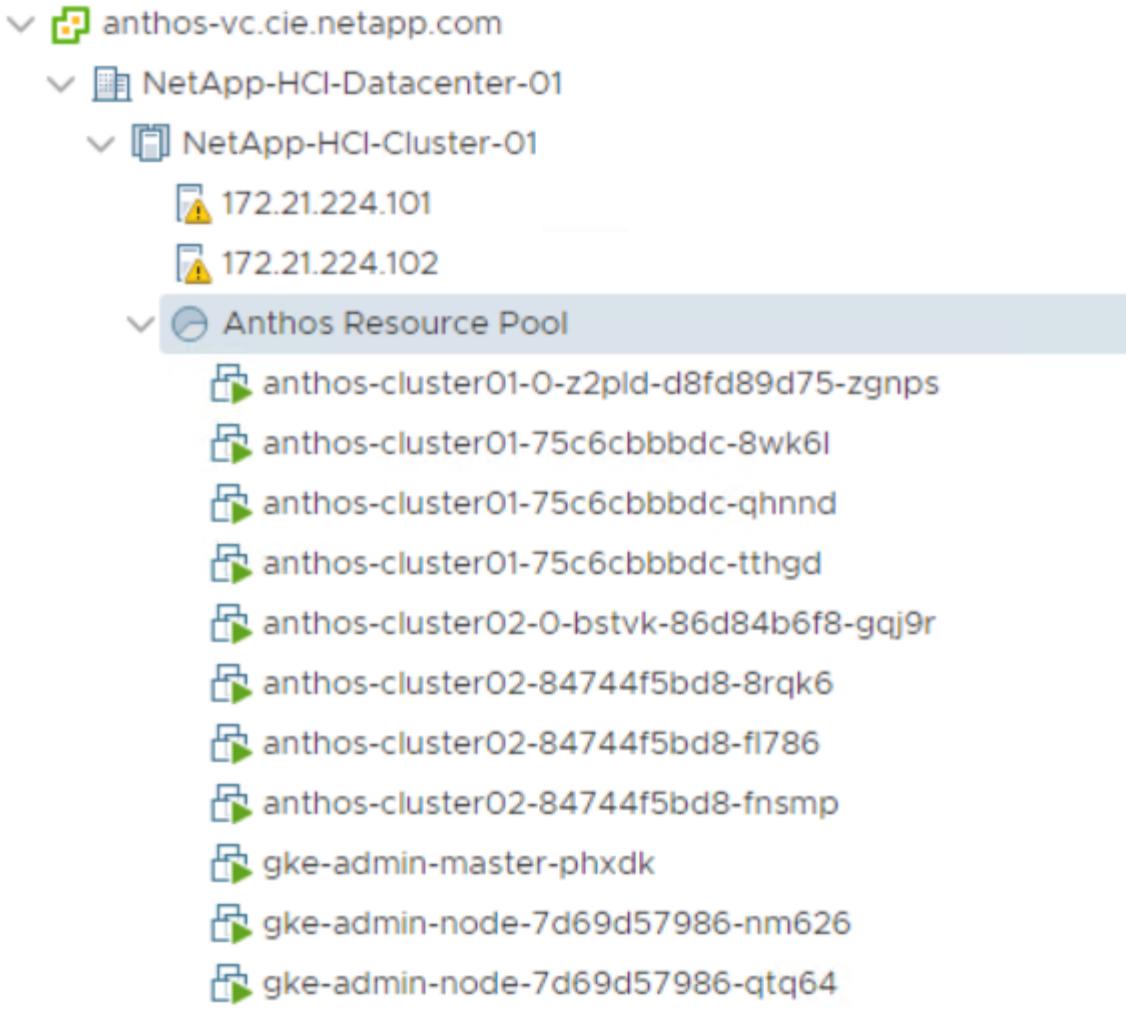
3. The process runs for several minutes and can be monitored on screen and in vCenter by watching the resource pool as the VMs populate. When complete, you should be able to see the `gke-admin` cluster (three nodes) and the first user cluster (four nodes).



During the deployment process, the standard out might display several messages about the current node not being available or not being ready. This is normal and happens when the control plane checks for machines that have not yet completed deployment or received DHCP addresses.



When using DHCP, if a deployment fails because nodes cannot be reached, there might not be enough available addresses in the pool. Leases for previously failed deployments might need to be cleared manually to allow for additional deployment attempts.



4. You can access and execute commands against the user cluster that has been deployed using the `kubectl` command line tool and the `kubeconfig` file generated by the process (stored in the working directory).

```
ubuntu@Anthos-Admin-Workstation:~$ kubectl get nodes --kubeconfig anthos-cluster01-kubeconfig
NAME                  STATUS  ROLES   AGE      VERSION
anthos-cluster01-75c6cbbbdc-8wk6l  Ready   <none>  149m    v1.13.7-gke.20
anthos-cluster01-75c6cbbbdc-qhnnd  Ready   <none>  149m    v1.13.7-gke.20
anthos-cluster01-75c6cbbbdc-tthgd  Ready   <none>  149m    v1.13.7-gke.20
```

9. Deploy Additional User Clusters: NetApp HCI with Anthos

With Anthos, organizations can scale their environments to incorporate multiple user clusters and segregate workloads between teams. A single admin cluster can

support up to five user clusters, and each user cluster can support up to twenty-five nodes.

To add additional user clusters to your deployment, complete the following steps:

1. Copy the `config.yaml` file to a new file named `anthos-cluster02-config.yaml`.

```
ubuntu@Anthos-Admin-Workstation:~$ cp config.yaml anthos-cluster02-config.yaml
```

2. Make the following edits to the newly created file:

1. Comment out the sections that refer to the existing admin cluster with (#).

2. When you get to the `usercluster` section, update the following fields:

1. Update the partition name under the `bigip` section.
2. Update the `controlplanvip` and `ingressvip` values under the `vip` section.
3. Update the `clusternamespace` value.

```
usercluster:  
  # In-Cluster vCenter configuration  
  vcenter:  
    # If specified it overwrites the network field in global vcenter  
    configuration  
    network: ""  
    # # The absolute or relative path to the yaml file to use for static IP  
    allocation.  
    # # Do not include if using DHCP  
    # ipblockfilepath: ""  
    # # Specify pre-defined nodeports if using "manual" load balancer mode  
    # manuallbspec:  
    #   ingresshttpnodeport: 30243  
    #   ingresshttpsnodeport: 30879  
    #   controlplanenodeport: 30562  
    #   addonsnodeport: 0  
    # Specify the already-existing partition and credentials to use with F5  
    bigip:  
      # To re-use credentials across clusters we recommend using YAML node  
      anchors.  
      # See https://yaml.org/spec/1.2/spec.html#id2785586  
      credentials:  
        address: "172.21.224.22"  
        username: "admin"  
        password: "NetApp!23"  
      partition: "Anthos-Cluster02-Part"  
      # # Optionally specify a pool name if using SNAT
```

```

# snatpoolname: ""
# The VIPs to use for load balancing
vips:
    # Used to connect to the Kubernetes API
    controlplanevip: "10.63.172.108"
    # Shared by all services for ingress traffic
    ingressvip: "10.63.172.109"
    # # Used for admin cluster addons (needed for multi cluster features). Must
be the same
    # # across clusters
    # addonsvip: ""

# A unique name for this cluster
clustername: "anthos-cluster02"
# User cluster master nodes must have either 1 or 3 replicas
masternode:
    cpus: 4
    memorymb: 8192
    # How many machines of this type to deploy
    replicas: 1
# The number of worker nodes to deploy and their size. Min. 2 replicas
workernode:
    cpus: 4
    memorymb: 8192
    # How many machines of this type to deploy
    replicas: 3
# The Kubernetes service CIDR range for the cluster
serviceiprange: 10.96.0.0/12
# The Kubernetes pod CIDR range for the cluster
podiprange: 192.168.0.0/16

```

3. Run the following command to check the config file again to verify that there are no syntax errors. Because you have removed the admin section, you must reference the **kubeconfig** file for the admin cluster named **kubeconfig** (found in the working directory).

```
ubuntu@Anthos-Admin-Workstation:~$ gkectl check-config --config anthos-cluster02-
config.yaml --kubeconfig kubeconfig
- Validation Category: Config Check
  - [SUCCESS] Config

- Validation Category: Docker Registry
  - [SUCCESS] gcr.io/gke-on-prem-release access

- Validation Category: vCenter
  - [SUCCESS] Credentials
  - [SUCCESS] Datacenter
  - [SUCCESS] Datastore
  - [FAILURE] Data Disk: vCenter data disk already exists
  - [SUCCESS] Resource Pool
  - [SUCCESS] Network

- Validation Category: F5 BIG-IP
  - [SUCCESS] Credentials
  - [SUCCESS] Partition

- Validation Category: Network Configuration
  - [SUCCESS] CIDR, VIP and static IP (availability and overlapping)

- Validation Category: VIPs
  - [SUCCESS] ping (availability)

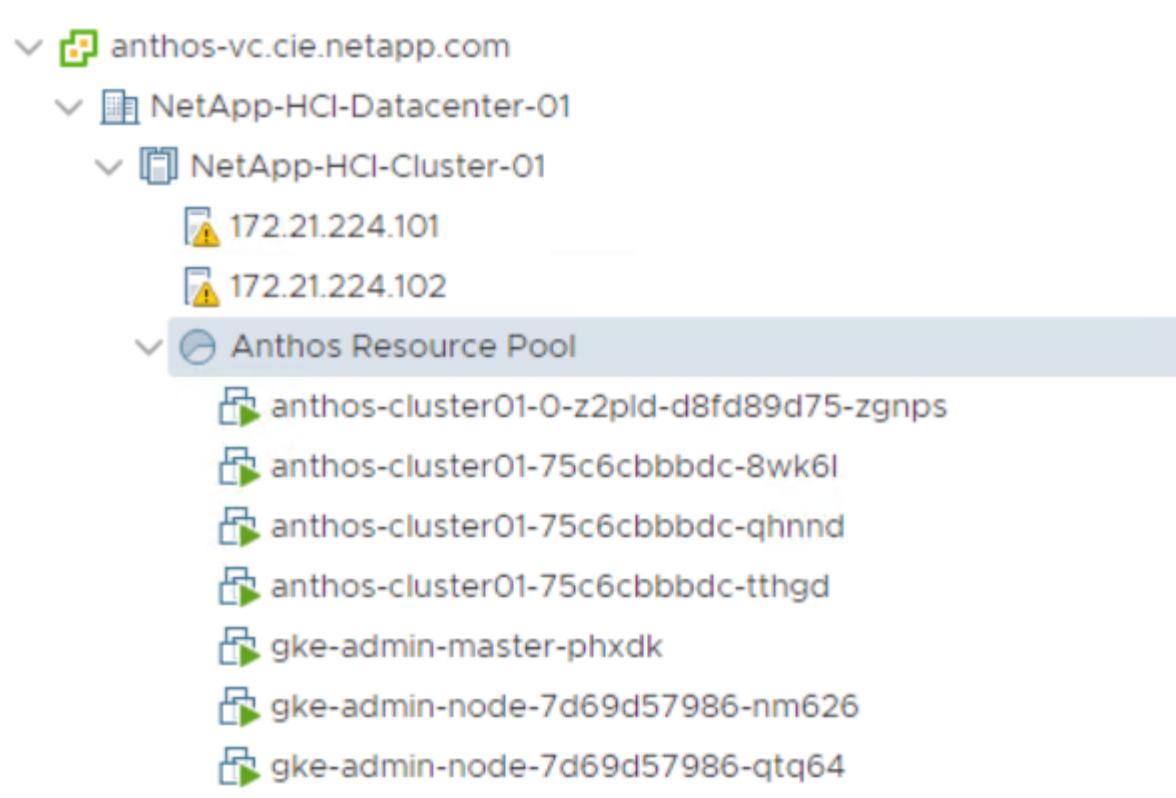
- Validation Category: Node IPs
  - [SUCCESS] ping (availability)
```

Some validations FAILED or SKIPPED. Check report above.

4. If all the checks succeed as expected, you can deploy this new user cluster in a manner very similar to the first cluster creation, referencing the **kubeconfig** file from the admin cluster.

```
ubuntu@Anthos-Admin-Workstation:~$ gkectl create cluster --config anthos-cluster02-
config.yaml --kubeconfig kubeconfig
```

5. As with the previous deployment, the process runs for several minutes and can be monitored on screen and in vCenter by watching the resource pool as the VMs populate. When complete, you should be able to see the new user cluster (four nodes).



6. You can access and execute commands against the deployed user cluster using the `kubectl` command line tool and the `kubeconfig` file generated by the process (stored in the working directory).

```
ubuntu@Anthos-Admin-Workstation:~$ kubectl get nodes --kubeconfig anthos-cluster02-kubeconfig
NAME                  STATUS  ROLES   AGE      VERSION
anthos-cluster02-84744f5bd8-8rqk6  Ready   <none>  9m16s  v1.13.7-gke.20
anthos-cluster02-84744f5bd8-f1786  Ready   <none>  9m28s  v1.13.7-gke.20
anthos-cluster02-84744f5bd8-fnsmp  Ready   <none>  9m21s  v1.13.7-gke.20
```

10. Enable Access to the Cluster with the GKE Console: NetApp HCI with Anthos

After clusters are deployed and registered with Google Cloud, they must be logged into with the Google Cloud console to be managed and to receive additional cluster details. The official procedure to gain access to Anthos user clusters after they are deployed is detailed [here](#).



The project and the specific user must be whitelisted to access on-premises clusters in the Google Cloud console and use Anthos on VMware services. If you are unable to see the clusters after they are deployed, you might need to open a support ticket with Google.

Figure 1. Non-whitelisted view.

Figure 2. View of clusters.

To enable access to your user clusters using the GKE console, complete the following steps:

1. Create a `node-reader.yaml` file that gives you the ability to access the cluster.

```
kind: clusterrole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: node-reader
rules:
- apiGroups: []
  resources: ["nodes"]
  verbs: ["get", "list", "watch"]
```

2. Apply this file to the cluster that you want to log into with the `kubectl` command.

```
ubuntu@Anthos-Admin-Workstation:~$ kubectl apply -f node-reader.yaml --kubeconfig
anthos-cluster01-kubeconfig
clusterrole.rbac.authorization.k8s.io/node-reader created
```

3. Create a Kubernetes service account (KSA) that you can use to login. Name this account after the user that uses this account to log into the cluster.

```
ubuntu@Anthos-Admin-Workstation:~$ kubectl create serviceaccount netapp-user
--kubeconfig anthos-cluster01-kubeconfig
serviceaccount/netapp-user created
```

4. Create cluster role-binding resources to bind both the view and newly created node-reader roles to the newly created KSA.

```
ubuntu@Anthos-Admin-Workstation:~$ kubectl create clusterrolebinding netapp-user-view  
--clusterrole view --serviceaccount default:netapp-user --kubeconfig anthos-cluster01-  
kubeconfig  
clusterrolebinding.rbac.authorization.k8s.io/netapp-user-view created  
ubuntu@Anthos-Admin-Workstation:~$ kubectl create clusterrolebinding netapp-user-node-  
reader --clusterrole node-reader -  
--serviceaccount default:netapp-user --kubeconfig anthos-cluster01-kubeconfig  
clusterrolebinding.rbac.authorization.k8s.io/netapp-user-node-reader created
```

5. If you need to extend permissions further, you can grant the KSA user a role with cluster admin permissions in a similar manner.

```
ubuntu@Anthos-Admin-Workstation:~$ kubectl create clusterrolebinding netapp-user-admin  
--clusterrole cluster-admin --serviceaccount default:netapp-user --kubeconfig anthos-  
cluster01-kubeconfig  
clusterrolebinding.rbac.authorization.k8s.io/netapp-user-admin created
```

6. With the KSA account created and assigned with correct permissions, you can create a bearer token to allow access with the GKE Console. To do so, set a system variable for the secret name, and pass that variable through a `kubectl` command to generate the token.

```
ubuntu@Anthos-Admin-Workstation:~$ SECRET_NAME=$(kubectl get serviceaccount netapp-  
user --kubeconfig anthos-cluster01-kubeconfig -o jsonpath='{$.secrets[0].name}')  
ubuntu@Anthos-Admin-Workstation:~$ kubectl get secret ${SECRET_NAME} --kubeconfig  
anthos-cluster01-kubeconfig -o jsonpath='{$.data.token}' | base64 -d  
eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlc5ldGVzL3NlcnpY2VhY2NvdW50Iiwia3ViZ  
XJuZXRlcyc5pb9zZXJ2aWN1YWNjb3VudC9uYW1lc3BhY2UiOijkZWZhWx0Iiwia3ViZXJuZXRlcyc5pb9zZXJ  
2aWN1YWNjb3VudC9zZWNyZXQubmFtZSI6Im5ldGFwcC11c2VyLXRva2VuLWJxd3piIiwia3ViZXJuZXRlcyc5pb  
y9zZXJ2aWN1YWNjb3VudC9zZXJ2aWN1LWFjY291bnQubmFtZSI6Im5ldGFwcC11c2VyIiwia3ViZXJuZXRlcyc5pb  
y9zZXJ2aWN1YWNjb3VudC9zZXJ2aWN1LWFjY291bnQudWlkIjoimIzzTFiZjQtMDE3NS0xMWVhLWEzMGutN  
mFiZmR1YjYwNDBmIwic3ViIjoic3lzdGvTOnNlcnpY2VhY2NvdW500mR1ZmF1bHQ6bmV0YXBwLXVzZXiifQ.  
YrHn4kYlb3gwxVKCLyo7p6J1f7mwwIgZqNw9eTvIk4PfyR4IJHxQwawnJ4T6R1jIFcbVSQwvWI1yGuTJ981AD  
dcwtFXHoEfMc0a6SIn40MVw1d5BGloaESn8150VCK3xES2DHAmLexFBqhVBgckZ0E4fZDvn4EhYvtFVpK1RbSy  
aE-  
DHD59P1bIgPdioiKREgbOddKdMn6XTVsui4V4tVKhktcdRNRAuw6cFDY1fPo13BFHr2aNBIe6lFLkUqvQN-  
9nMd63JGdHL4hfXu6PPDxc9By6Lg0W0nyaH4__gexy4uIa61fNLKV2SKe4_gAN41ff0CKe4Tq8sa6zMo-8g
```

7. With this token, you can visit the [Google Cloud Console](#) and log in to the cluster by clicking the login button and pasting in the token.

Log in to cluster

Choose the method you want to use for authentication to the cluster

Token

`Jxc9By6Lg0W0nyaH4_gexy4ula61fNLKV2SKe4_gAN41ffOCKe4Tq8sa6zMo-8g`

- Basic authentication
- Authenticate with Identity Provider configured for the cluster

[CLOSE](#) [LOGIN](#)

1. After login is complete, you see a green check mark next to the cluster name, and information is displayed about the physical environment. Clicking the cluster name displays more verbose information.

11. Install and Configure NetApp Trident Storage Provisioner: NetApp HCI with Anthos

Trident is a storage orchestrator for containers. With Trident, microservices and containerized applications can take advantage of enterprise-class storage services provided by the full NetApp portfolio of storage systems for persistent storage mounts. Depending on an application's requirements, Trident dynamically provisions storage for ONTAP-based products such as NetApp AFF and FAS systems and Element storage systems like NetApp SolidFire® and NetApp HCI.

To install Trident on the deployed user cluster and provision a persistent volume, complete the following steps:

1. Download the installation archive to the admin workstation and extract the contents. The current version of Trident is 19.10, which can be downloaded [here](#).

```
ubuntu@Anthos-Admin-Workstation:~$ wget
https://github.com/NetApp/trident/releases/download/v19.10.0/trident-installer-
19.10.0.tar.gz
--2019-11-07 16:45:33--
https://github.com/NetApp/trident/releases/download/v19.10.0/trident-installer-
19.10.0.tar.gz
Resolving github.com (github.com)... 140.82.118.4
Connecting to github.com (github.com)|140.82.118.4|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://github-production-release-asset-
2e65be.s3.amazonaws.com/77179634/4d3b5900-
Resolving github-production-release-asset-2e65be.s3.amazonaws.com (github-production-
release-asset-2e65be.s3.amazonaws.com)... 52.216.81.8
Connecting to github-production-release-asset-2e65be.s3.amazonaws.com (github-
production-release-asset-2e65be.s3.amazonaws.com)|52.216.81.8|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 68903585 (66M) [application/octet-stream]
Saving to: "trident-installer-19.10.0.tar.gz"

trident-installer-19.10.0.tar 100%[=====] 65.71M 53.8MB/s in 1.2s

2019-11-07 16:45:35 (53.8 MB/s) - "trident-installer-19.10.0.tar.gz" saved
[68903585/68903585]

ubuntu@Anthos-Admin-Workstation:~$ tar -xf trident-installer-19.10.0.tar.gz
```

2. First set the location of the user cluster's **kubeconfig** file as an environment variable so that you don't have to reference it, because Trident has no option to pass this file.

```
ubuntu@Anthos-Admin-Workstation:~$ export KUBECONFIG=~/anthos-cluster01-kubeconfig
```

3. Navigate to the Trident directory and execute the `tridentctl` tool to install Trident to your cluster. NetApp recommends installing Trident into its own namespace within the cluster. You can then verify that the install finished correctly.

```

ubuntu@Anthos-Admin-Workstation:~$ cd trident-installer --csi
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ ./tridentctl install -n trident
INFO Starting Trident installation.           namespace=trident
INFO Created namespace.                      namespace=trident
INFO Created service account.
INFO Created cluster role.
INFO Created cluster role binding.
INFO Created custom resource definitions.    namespace=trident
INFO Added finalizers to custom resource definitions.
INFO Created Trident deployment.
INFO Waiting for Trident pod to start.
INFO Trident pod started.                   namespace=trident pod=trident-
79c76ff764-77sgl
INFO Waiting for Trident REST interface.
INFO Trident REST interface is up.          version=19.10.0
INFO Trident installation succeeded.
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ ./tridentctl version -n trident
+-----+-----+
| SERVER VERSION | CLIENT VERSION |
+-----+-----+
| 19.10.0        | 19.10.0       |
+-----+-----+

```

4. The next step in enabling Trident integration with the NetApp HCI solution and Anthos is to create a backend that enables communication with the storage system. There are sample backend files available in the downloaded installation archive in the `sample-input` folder. Copy the `backend-solidfire.json` to your working directory and edit it to provide information detailing the storage system environment.

```

ubuntu@Anthos-Admin-Workstation:~/trident-installer$ cp sample-input/backend-
solidfire.json ./
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ vi backend-solidfire.json

```

5. Edit the user, password, and MVIP value on the EndPoint line.
6. Edit the SVIP value.

```
{
  "version": 1,
  "storageDriverName": "solidfire-san",
  "Endpoint": "https://trident:password@172.21.224.150/json-rpc/8.0",
  "SVIP": "10.63.172.100:3260",
  "TenantName": "trident",
  "Types": [{"Type": "Bronze", "Qos": {"minIOPS": 1000, "maxIOPS": 2000, "burstIOPS": 4000}, {"Type": "Silver", "Qos": {"minIOPS": 4000, "maxIOPS": 6000, "burstIOPS": 8000}, {"Type": "Gold", "Qos": {"minIOPS": 6000, "maxIOPS": 8000, "burstIOPS": 10000}}}]
}
```

- With this back-end file in place, run the following command to create your first backend.

```
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ ./tridentctl -n trident create
backend -f backend.json
+-----+-----+-----+-----+
+-----+
|   NAME          | STORAGE DRIVER |           UUID           | STATE |
VOLUMES |
+-----+-----+-----+-----+
+-----+
| solidfire-backend | solidfire-san | a5f9e159-c8f4-4340-a13a-c615fef0f433 | online |
0 |
+-----+-----+-----+-----+
+-----+
```

- With the backend created, you must next create a storage class. Just as with the backend, there is a sample storage class file that can be edited for the environment available in the sample-inputs folder. Copy it to the working directory and make necessary edits to reflect the backend created.

```
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ cp sample-input/storage-class-
csi.yaml.templ ./storage-class-basic.yaml
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ vi storage-class-basic.yaml
```

- The only edit that must be made to this file is to define the `backendType` value to the name of the storage driver from the newly created backend. Also note the `name-field` value, which must be referenced in a later step.

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: basic-csi
  provisioner: csi.trident.netapp.io
parameters:
  backendType: "solidfire-san"
```

- Run the `kubectl` command to create the storage class.

```
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ kubectl create -f sample-input/storage-class-basic.yaml
```

- With the storage class created, you must then create the first persistent volume claim (PVC). There is a sample `pvc-basic.yaml` file that can be used to perform this action located in `sample-inputs` as well. The only edit that must be made to this file is ensuring that the `storageClassName` field matches the one just created.

```
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ vi sample-input/pvc-basic.yaml
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: basic
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  storageClassName: basic-csi
```

- Create the PVC by issuing the `kubectl` command, Creation can take some time depending on the size of the backing volume being created, so you can watch the process as it completes.

```
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ kubectl create -f sample-input/pvc-basic.yaml
```

```
ubuntu@Anthos-Admin-Workstation:~/trident-installer$ kubectl get pvc --watch
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES
STORAGECLASS	AGE			
basic	Pending			
basic	1s			
basic	Pending	pvc-2azg0d2c-b13e-12e6-8d5f-5342040d22bf	0	
basic	5s			
basic	Bound	pvc-2azg0d2c-b13e-12e6-8d5f-5342040d22bf	1Gi	RWO
basic	7s			

Copyright Information

Copyright © 2020 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.